

The background image shows an aerial view of the University of Salerno's campus in Italy. The campus features modern buildings with unique architectural designs, including a large dome and various geometric shapes. A large open plaza with a grid pattern is visible, populated by many people. In the distance, a range of green mountains under a clear blue sky provides a scenic backdrop.

FPGA

High Performance Computing, Summer 2021



Biagio Cosenza
Department of Computer Science
University of Salerno
bcosenza@unisa.it


Outline

- FPGA
 - spatial architecture
 - pipeline
 - data flow
 - pipelining, queues, pipes




Von Neumann Bottleneck

- Load-store “Von Neumann” architecture
 - based on Instruction Set Architecture (ISA)
 - e.g., CPUs and GPUs
 - executes a different instruction from the program in each clock cycle
- Hardware reuse
 - hardware architecture that can run different instructions at different times
 - e.g., the memory load unit feeding an addition is probably the same memory load unit used to feed a subtraction
 - e.g., the same arithmetic unit is probably used to execute both the addition and subtraction instructions





Spatial Architecture

- Spatial implementations of a program conceptually take the entire program as a whole and lay it down at once on the device
 - different regions of the device implement different instructions in the program
 - each instruction receives its own dedicated hardware that can execute simultaneously (same clock cycle) as the hardware implementing the other instructions
 - opposite perspective from sharing hardware between instructions over time (e.g., ISA)
- Static **dataflow** architecture
 - No longer need: instruction fetch, decode unit, program counter, register file
 - Dataflow: connect the output of one instruction to the input of the next



Von Neumann vs Spatial Architectures



- Von Neumann arch.
- Energy break down for a single add
 - I-Cache access: 25pJ
 - register file access: 6pJ
 - control: 29 pJ
 - total energy per instruction 70pJ
- Static Dataflow
- Energy per operation: 1-3pJ

$$x = a + b$$



Overcoming the Von Neumann Bottleneck

- Control is the bottleneck of Von Neumann architectures
- SIMD vector on CPUs and SIMT on GPUs
 - reduce the control overhead
 - but still need extra energy for registers
- FPGA reduce overhead for registers as well as control



FPGA Fitting

- Each instruction occupies some percentage of the spatial area of the device
 - What happens if the program requires more than 100% of the area?
- Small programs
 - if a program uses most of the area on the FPGA and there is sufficient work to keep all of the hardware busy every clock cycle
 - then executing a program can be incredibly efficient because of the extreme parallelism
- Large programs
 - to be tuned and restructured to fit on a device
 - **resource sharing** features of compilers can help to address this
 - but usually with some degradation in performance that reduces the benefit of using an FPGA
 - ISA-based architecture are very efficient resource sharing





FPGA Resource Sharing

- Taken to the extreme, resource sharing solutions on an FPGA lead to an architecture that looks like an ISA-based accelerator
 - but that is built in reconfigurable logic instead being optimized in fixed silicon
- The reconfigurable logic leads to overhead relative to a fixed silicon design
 - therefore, FPGAs are not typically chosen as ways to implement ISAs
- FPGAs are of prime benefit
 - when an application is able to utilize the resources to implement efficient data flow algorithms




Pipeline Parallelism

- No pipeline
 - quite inefficient
 - most of the time
 - the hw is only doing useful work a small percentage of the time
 - operation such as the multiply is either waiting for new data from the load or holding its output so that operations later in the chain can use its result
- With pipeline
 - execution of a single program is spread across many cc
 - achieved by inserting registers between some operations
 - each register holds a binary value for the duration of a cc
 - holds the result of an operation's output so that the next operation in the chain can operate on that held value
 - the previous operation is free to work on a different computation without impacting the input to following operations



Pipeline with Control Flow

- On ISA-based architectures
 - control flow (e.g., if/else structures) leaves some parts of the program inactive a certain percentage of the clock cycles
- FPGA compilers typically combine hardware from both sides of a branch, where possible
 - to minimize wasted spatial area and to maximize compute efficiency during control flow divergence
 - flow divergence much less expensive




Scalar Data Flow

- Intermediate data between **operations** not only stays on-chip
 - is not stored to external memory
- Intermediate data between each **pipeline stage** has dedicated storage registers
 - different from vector architectures where multiple computations are executed as lanes of a shared vector instruction
- FPGAs are good for algorithms where
 - data dependences across units of work (such as work-items) can't be broken
 - fine-grained communication must occur



FPGA Chip Space

- Each kernel consumes **chip area** or **space** on the device
- Multiple kernels in the same FPGA binary
 - kernels can run concurrently
 - independent forward progress between kernels



FPGA Occupancy

- If there isn't enough work to occupy most of the pipeline stages most of the time, then efficiency will be low
- We'll call the average utilization of pipeline stages over time **occupancy** of the pipeline
 - different from the definition of occupancy used when optimizing other architectures such as GPUs




FPGA Flexibility with Custom Operations

- FPGAs are very efficient at
 - bitwise operations, integer math operations on custom data widths or types
 - e.g., a 33-bit integer multiplier or a 129-bit adder
 - common in image processing and machine learning




FPGA Low Latency I/O Streaming

- Some FPGA accelerator cards have network interfaces
 - make it possible to stream data directly into the device, process it, and then stream the result directly back to the network
 - good when
 - latency needs to be minimized
 - where processing through operating system network stacks is too slow or needs to be offloaded



FPGA Customized Memory System

- Memory systems built out of small blocks of on-chip memory
 - custom built for the specific portion of an algorithm or kernel using it
 - good for: applications that have atypical memory access patterns and structures
- FPGA memory systems are customized by the compiler for our specific code
 - GPU: avoid bank conflicts
 - FPGA: can change num of access ports per bank or have an unusual number of banks



FPGA: Performance Caveat

- FPGA devices differ from vendor to vendor or even from product generation to product generation
 - best practices for one device may not be best practices for a different device
 - to achieve optimal performance for a particular FPGA, always consult the vendor's documentation
- Most popular FPGA
 - Intel Stratix 10
 - 14 nm, 1 GhZ, 5.5M logic elements, TDP 125 W
 - Xilinx Virtex Ultrascale+
 - 14 nm, 600 MHz, 2.5 M logic elements, TDP 225W



FPGA Programming Models:

- Hardware Description Languages
 - Example: a counter in Verilog HDL
- High-level programming approaches
 - OpenCL / SYCL


```
module COUNTER #(  
    pWIDTH=8  
)  
(  
    input          iCLK,  
    input          iRESET,  
    output [pWIDTH-1:0] oCOUNTER  
,  
    reg [pWIDTH-1:0] rCOUNTER;  
    wire [pWIDTH-1:0] wNEXT_COUNTER;  
    assign wNEXT_COUNTER = rCOUNTER+1;  
    assign oCOUNTER = rCOUNTER;  
    always @ (posedge iCLK)  
    begin  
        if (iRESET) begin  
            rCOUNTER<=0;  
        end else begin  
            rCOUNTER<= wNEXT_COUNTER;  
        end  
    end  
endmodule
```



FPGA Programming Models: SYCL/DPC++


- Intel Toolchain
- FPGA binaries are embedded within the compiled DPC++ executable that run on the host
 - FPGA is automatically configured
- When we run a host program and submit the first kernel for execution on an FPGA, there might be a delay before the kernel begins executing, while the FPGA is programmed
 - resubmitting kernels for additional executions won't see the same delay because the kernel is already programmed to the device and ready to run

```
#include <CL/sycl.hpp>
#include <CL/sycl/intel/fpga_extensions.hpp>
using namespace sycl;
int main() {
    queue Q{ INTEL::fpga_selector{} };
    Q.submit([&] (handler &h) {
        h.parallel_for(1024, [=] (auto idx) {
            // ...
        });
    });
    return 0;
}
```




FPGA with SYCL/DPC++

- **Emulator**: acts as if it was an FPGA, including supporting relevant extensions and emulating the execution model
 - but runs on the host processor
 - not as fast as on actual FPGA
- queue Q{ INTEL::fpga_emulator_selector{} };
- **Static reports**: report on the FPGA structures created by the compiler and on bottlenecks identified by the compiler



Pipelining and Queues


- Goal of **pipelining**: enable multiple elements of data to be processed at different stages of the pipeline, simultaneously
- Example of inefficient pipeline
 - pipeline stages are mostly unused if processing only a single element of work



Pipelining and Queues

- Keep the pipeline stages better occupied

- queue of un-started work waiting before the first stage, which feeds the pipeline
- each clock cycle, the pipeline can consume and start one more element of work from the queue
- after some initial startup cycles, each stage of the pipeline is occupied





FPGA Queue Feeding in SYCL

- Two methods
 - ND-range kernels
 - Loops
- Optimization must carefully balance between them



FPGA Queue from ND Range

- ND-range execution model
 - work-items do not frequently communicate with each other in most applications



FPGA Backward Communication

- Loop-carried data dependence


```
int state = 0;  
for (int i=0; i < size; i++) {  
    state = generate_random_number(state);  
    output[i] = state;  
}
```

- FPGA can very efficiently communicate results backward in the pipeline

- to work at an earlier stage in the pipeline

- Approaches




1. loops
2. intra-kernel pipes with ND-range kernels



FPGA Queue fed by Loops



■ Loop pipelining

```
h.single_task([=] () {
    int state = seed;
    for (int i=0; i < size; i++) {
        state = generate_incremental_random_number(state);
        output[i] = state;
    }
});
```




FPGA Loops and Occupancy

- Pipeline can initiate a new loop iteration over N clock cycles
 - called **initiation interval**




Pipes

- Spatial architectures relies on a first-in first-out (FIFO) buffer
 - 1. implicit control information carried alongside the data, signals tell us whether the FIFO is empty or full and can be useful when decomposing a problem into independent pieces
 - 2. FIFOs have storage capacity: easier to achieve performance in the presence of dynamic behaviors such as highly variable latencies when accessing memory
- Used for modular design



Pipes

■ Four types of pipe connectivity in SYCL/DPC++



■ Example

1. ND range
2. Single task with a loop

```
using my_pipe = pipe<class some_pipe, int>;
// ND-range kernel
Q.submit([&] (handler& h) {
    auto A = accessor(B_in, h);
    h.parallel_for(count, [=] (auto idx) {
        my_pipe::write( A[idx] );
    });
});
// Single_task kernel
Q.submit([&] (handler& h) {
    auto A = accessor(B_out, h);
    h.single_task([=] () {
        for (int i=0; i < count; i++) {
            A[i] = my_pipe::read();
        }
    });
});
```



FPGA Custom Memory Systems

- Static coalescing
 - memory accesses into a smaller number of wider accesses, where it can
 - reduces the complexity of a memory system
 - in terms of numbers of load or store units in the pipeline, ports on the memory system, the size and complexity of arbitration networks, and other memory system details
- Memory access style
 - load or store units for memory accesses
 - tailored to the memory technology being accessed
 - e.g., on-chip vs. DDR vs. HBM
 - tailored to the access pattern inferred from the source code
 - e.g., streaming, dynamically coalesced/widened, or likely to benefit from a cache of a specific size
- Memory system structure
 - memory systems (both on- and off-chip) can have banked structures
 - many controls and mode modifications that can be used to control these structures and to tune specific aspects of the spatial implementation



FPGA Summary: Basic Elements

- Lookup table
- Math engine
 - for common math operations such as addition or multiplication of single-precision floating point numbers, FPGAs have specialized hardware to make those operations very efficient
- On-chip memory
 - registers that are used to pipeline between operations and some other purposes
 - block memories that provide small random-access memories spread across the device
- Interfaces to off-chip hardware
- Routing fabric between all of the other elements



Lab

- Project selection, to define
 1. Topic
 2. Target hardware
 3. Baseline
 4. Tasks
- Teaching assistant
 - Alessia Antelmi for projects on hypergraphs
 - Kaijie Fan for projects on energy
 - Majid Salimibeni for GPU
 - I will assist for other topics (vectorization, ect)

