

Optimizing Compilers

High Performance Computing, Summer 2021



Biagio Cosenza
Department of Computer Science
University of Salerno
bcosenza@unisa.it

Outline

- Introduction to compilers
- Domain Specific Languages
- Intermediated Representations
- LLVM



What is a Compiler?

"a computer program that translates a program written in a high-level language into another language, usually machine language"

source: <http://dictionary.reference.com>

- German translations (from <http://dict.leo.org>)
 - Der Kompilierer
 - Der Übersetzer
 - Very interesting translation
 - Literally: translator



What is a Compiler?



- What is a **compiler**?

a program that accepts as input a program text in a certain language and produces as output a program text in another language, while preserving the meaning of that text [Grune et al., 2000]

a program that reads a program written in one language (source language) and translates it into an equivalent program in another language (target language) [AHO]

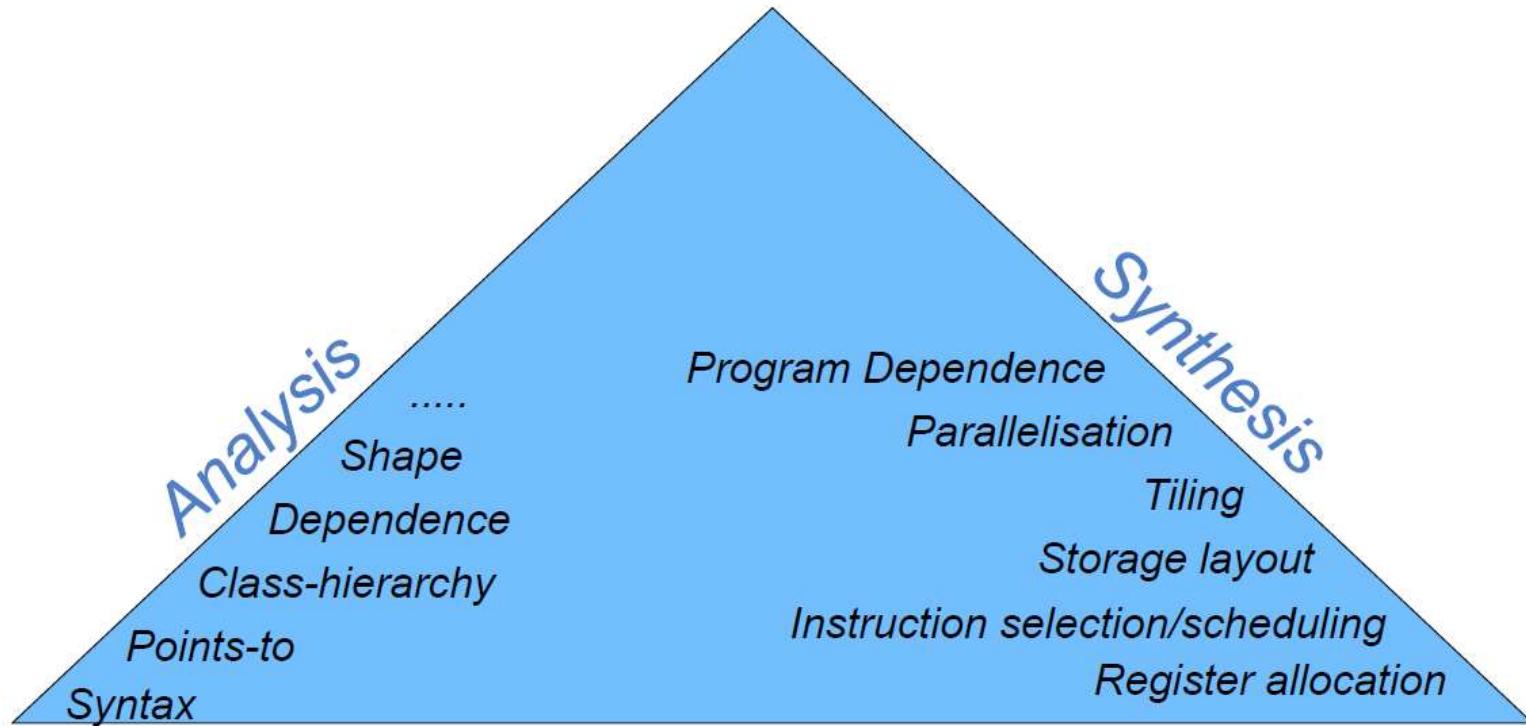
- key: ability to extract properties of a source program (**analysis**) and transform it to construct a target program (**synthesis**)

- What is an **interpreter**?

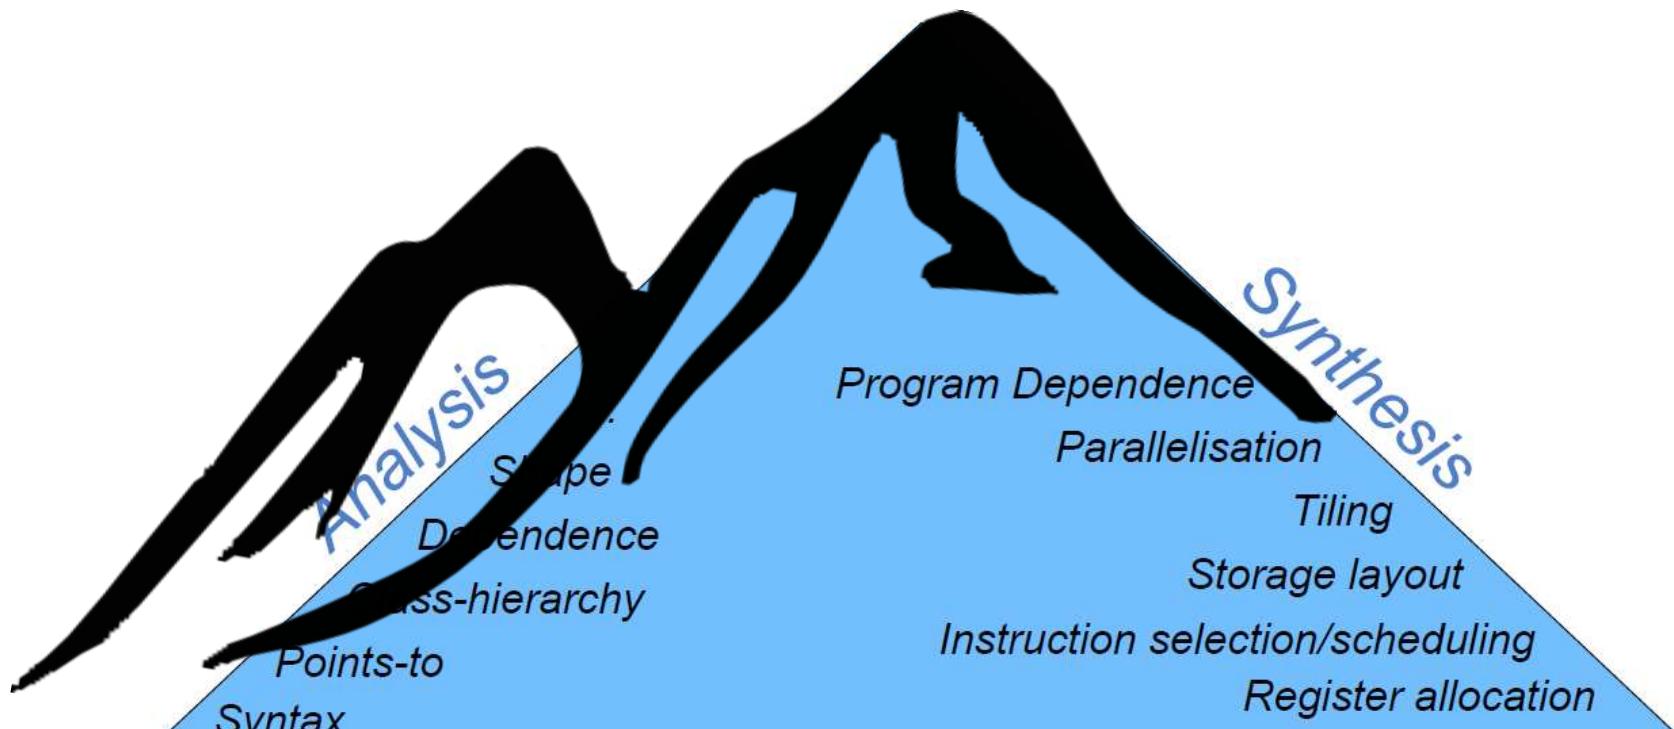
- a program that reads a source program and produces the results of executing this source
- an interpreter directly executes, i.e. performs, instructions written in a programming language, without previously compiling them into a machine language program



Analysis & Synthesis



Analysis & Synthesis



Courtesy of Paul Kelly, Imperial College London



Example 1: Source-to-source Compilers

- Also called **transpiler** or **transcompiler**
 - e.g., C-to-C
- Code transformation at source level
 - e.g., automatic parallelization, data layout transformations, ...
- High-level intermediate representation
- Examples
 - Rose, Insieme, Pluto, Cetus, ...



Example 2: Java JIT Compiler

- Java compiler
 - the output is a class file (.class)
 - javac (Oracle), gcj (GNU Compiler for Java), ECJ (Eclipse for Java), ...
 - platform-neutral Java bytecode
 - there are also compilers that emit optimized native machine code for a particular hardware/operating system combination
- Most Java-to-bytecode compilers do little optimization, leaving this to the JRE (Java Runtime) at runtime
- Just-in-time (JIT) compilation
 - the Java virtual machine (JVM) loads the class files and either interprets the bytecode or just-in-time compiles it to machine code and then possibly optimizes it using dynamic compilation.
 - interaction between JVM and Java compilers specified in JSR 199



Example 3: Domain-Specific Languages

- Input is a domain-specific language (DSL)
 - a language with domain-specific construct and constraints
- Assumptions (and restriction) on the input
 - analysis simpler
 - optimization is relatively easier
 - Domain-specific optimization
- Examples
 - OpenGL Shading Language
 - Halide for image processing
 - SQL for database

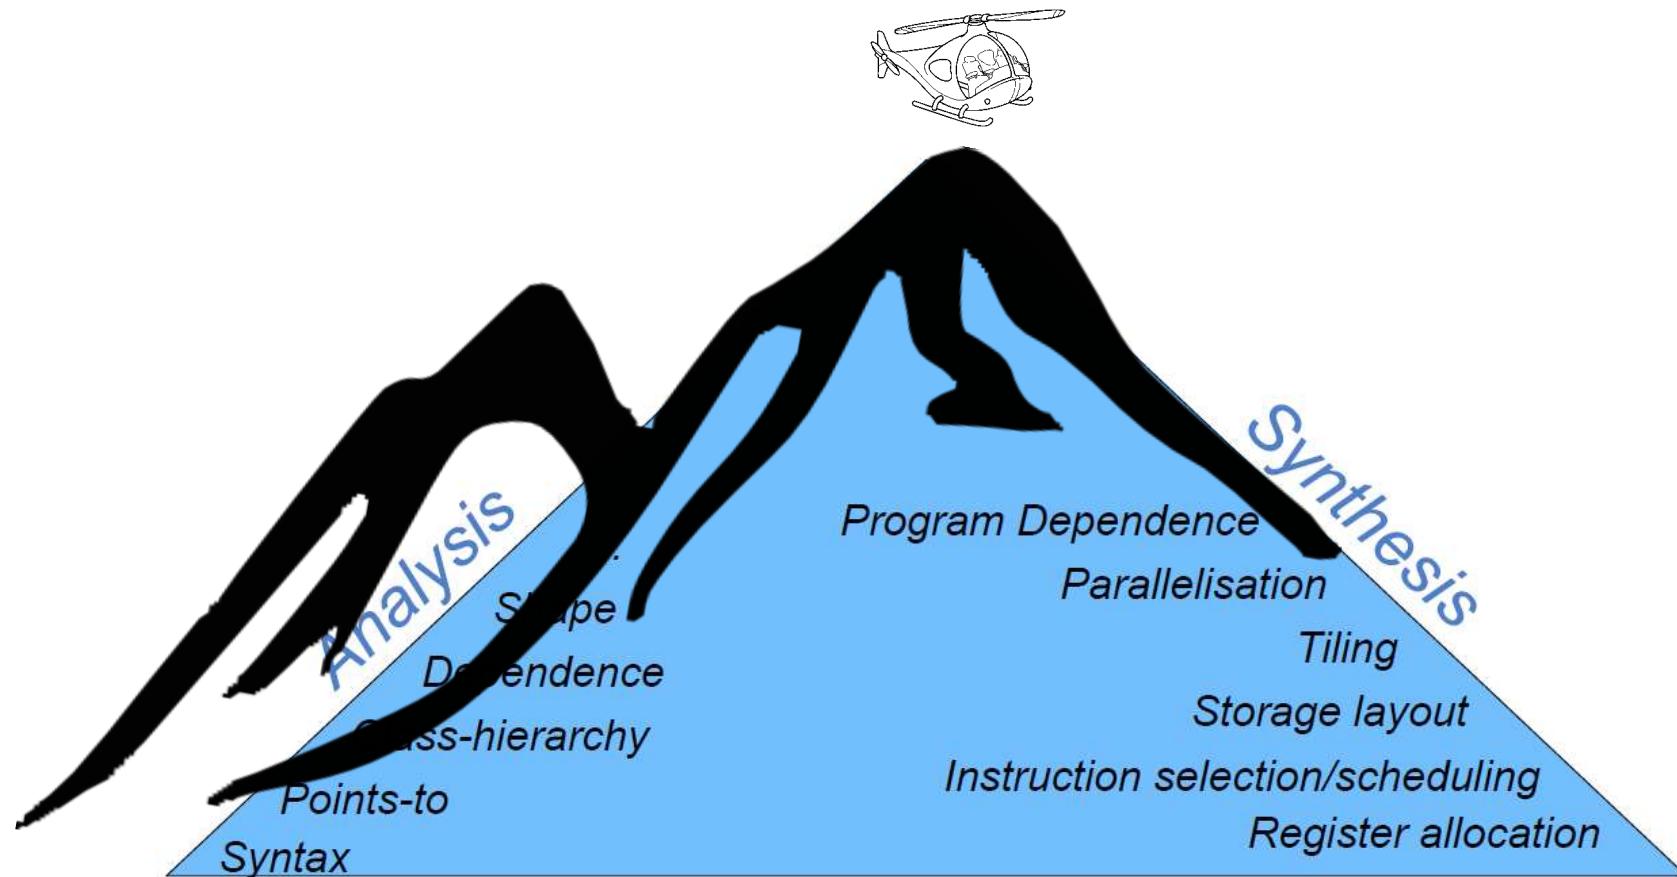
Example: simple GLSL fragment shader

```
varying vec3 N;
varying vec3 v;

void main(void) {
    vec3 L = normalize(gl_LightSource[0].position.xyz-v);
    vec4 Idiff = gl_FrontLightProduct[0].diffuse * max(dot(N,L), 0.0);
    Idiff = clamp(Idiff, 0.0, 1.0);
    gl_FragColor = Idiff;
}
```



DSL: - Analysis, + Synthesis



Courtesy of Paul Kelly, Imperial College London



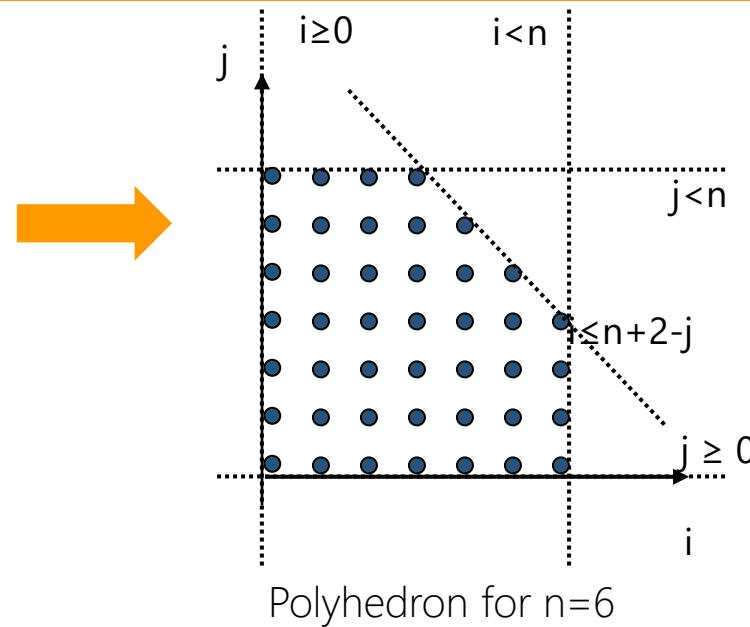
Example 4: Parallelizing Compilers

- Input is sequential code
- Output is parallel code, i.e., expose some kind of parallelism
- Typically, parallelism is extracted from loops
 - advanced analysis, e.g., using the **polyhedral model**
 - `#pragma` notations to help the compiler job
- Sometime parallelizing compilers enhance parallelization
 - E.g., from shared memory parallel code to distributed or heterogeneous systems
- Form of parallelism
 - automatic vectorization (SIMD instructions), e.g., by `gcc`, `llvm` and `icc`
 - multi-threading (`pthread`), e.g., by Rose, Pluto, Insieme, LLVM-Polly
 - distributed memory (typically MPI)



Automatic Parallelization with the Polyhedral Model

```
for(int i=0; i<n; i++)
    for(int j=0; j<n; j++)
        if(i <= n+2-j)
(s)            b[j] = b[j] + a[i];
```



i j n constant

$$\begin{bmatrix} 1 & 0 & 0 & 0 \\ -1 & 0 & 1 & -1 \\ 0 & 1 & 0 & 0 \\ 0 & -1 & 1 & -1 \\ -1 & -1 & -1 & -2 \end{bmatrix} \begin{pmatrix} i \\ j \\ n \\ 1 \end{pmatrix} \geq \vec{0}$$

Iteration domain with homogenous coordinates

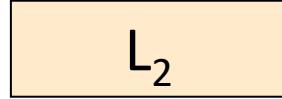
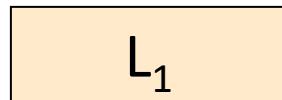
$$\begin{bmatrix} 1 & 0 \\ -1 & 0 \\ 0 & 1 \\ 0 & -1 \\ -1 & -1 \end{bmatrix} \begin{pmatrix} i \\ j \\ n \\ 0 \\ 1 \end{pmatrix} + \begin{pmatrix} 0 \\ n-1 \\ 0 \\ n-1 \\ -n-2 \end{pmatrix} \geq \vec{0}$$

Iteration domain of S

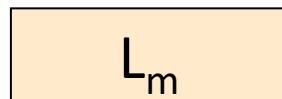


From Languages to Target Architectures

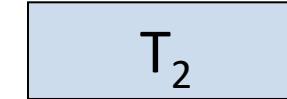
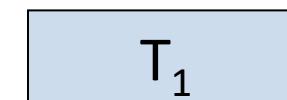
Programming Languages



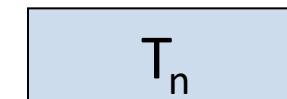
⋮



Target Architecture



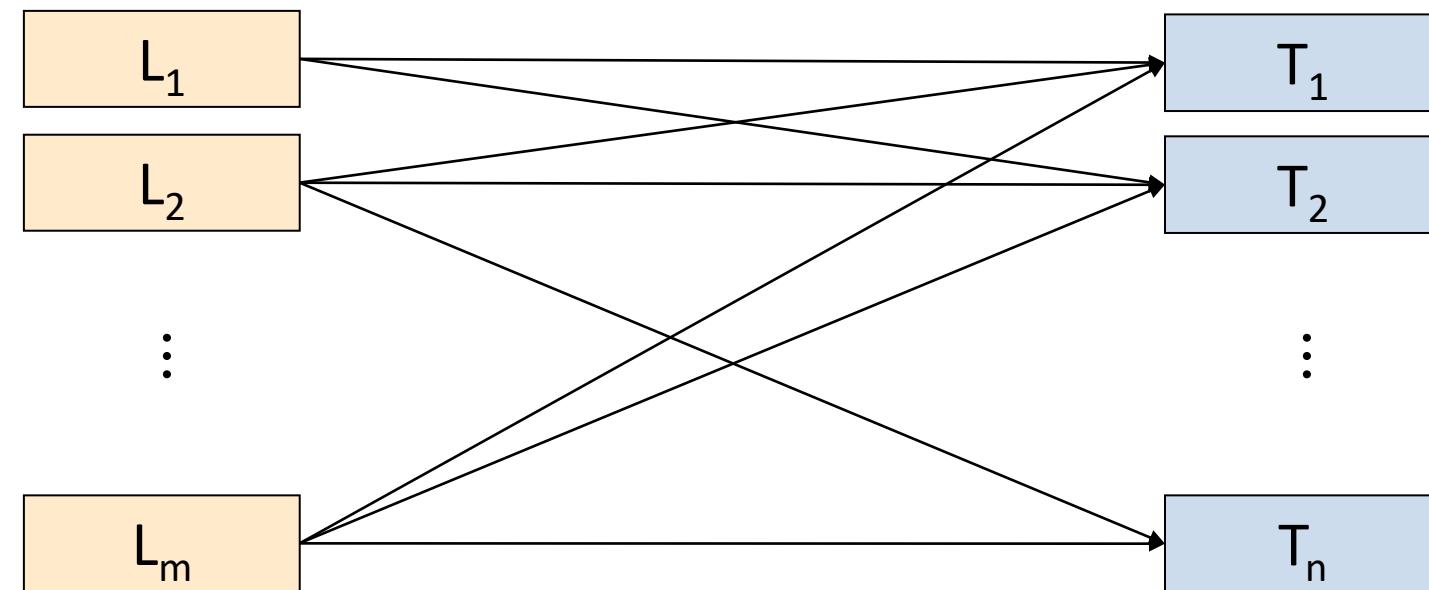
⋮



From Languages to Target Architectures

Programming Languages

Target Architecture



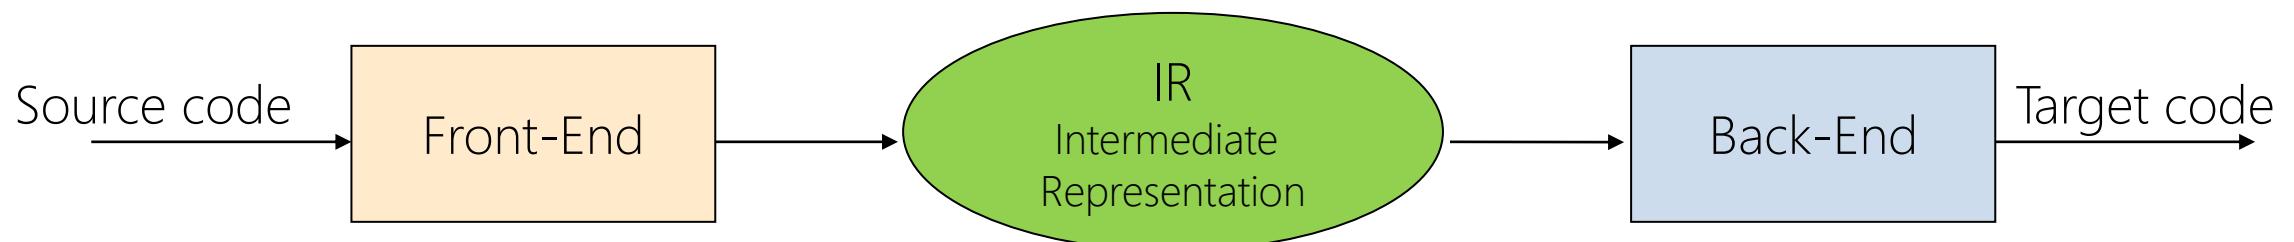
General Structure of a Compiler

Front-end performs the analysis of the source language

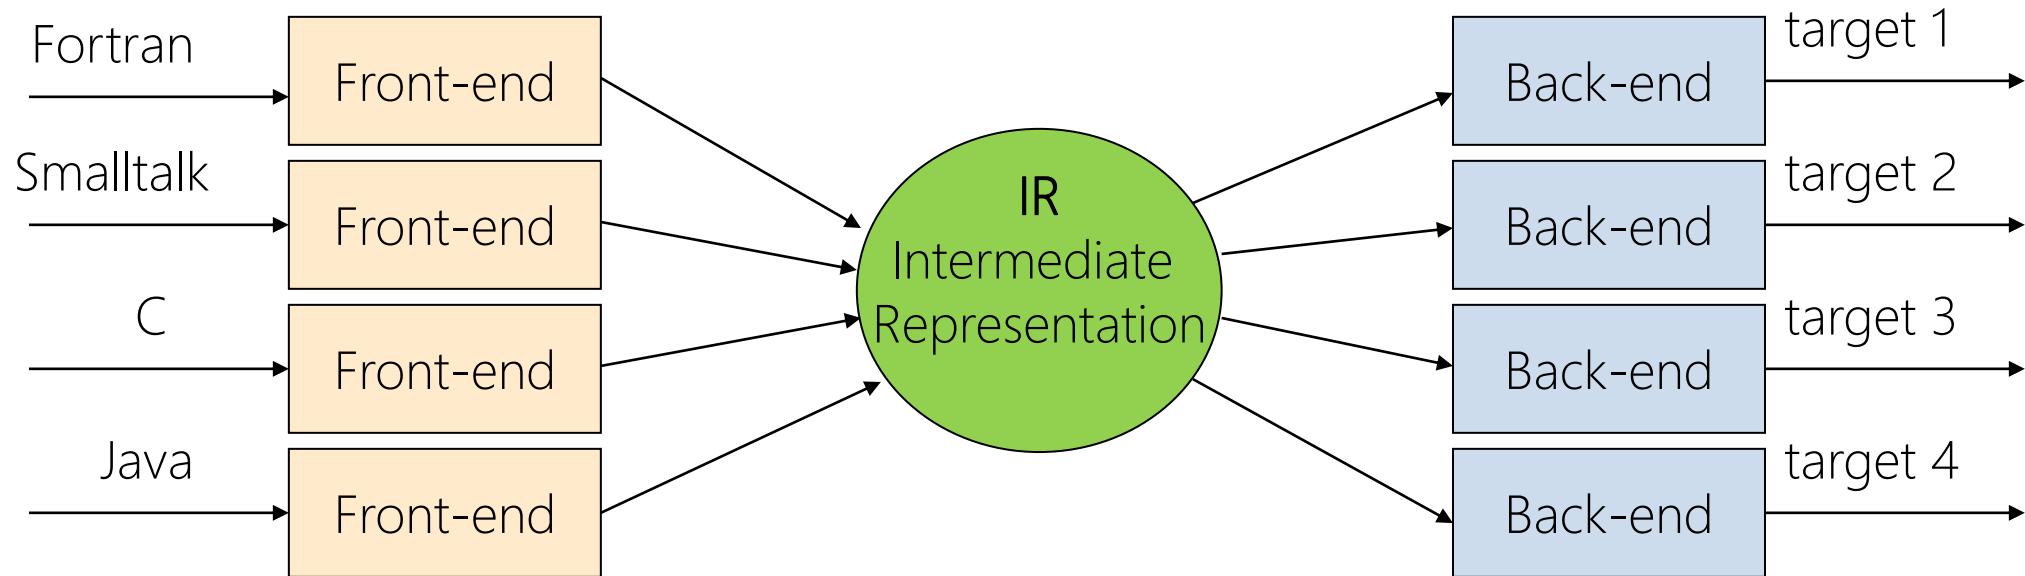
- recognizes legal and illegal programs and reports errors
- understands the input program and collects its semantics in an **Intermediate Representation (IR)**
- produces IR and shapes the code for the back-end

Back-end does the target language synthesis

- chooses instructions to implement each IR operation
- translates IR into target code
- needs to conform with system interfaces
- automation has been less successful



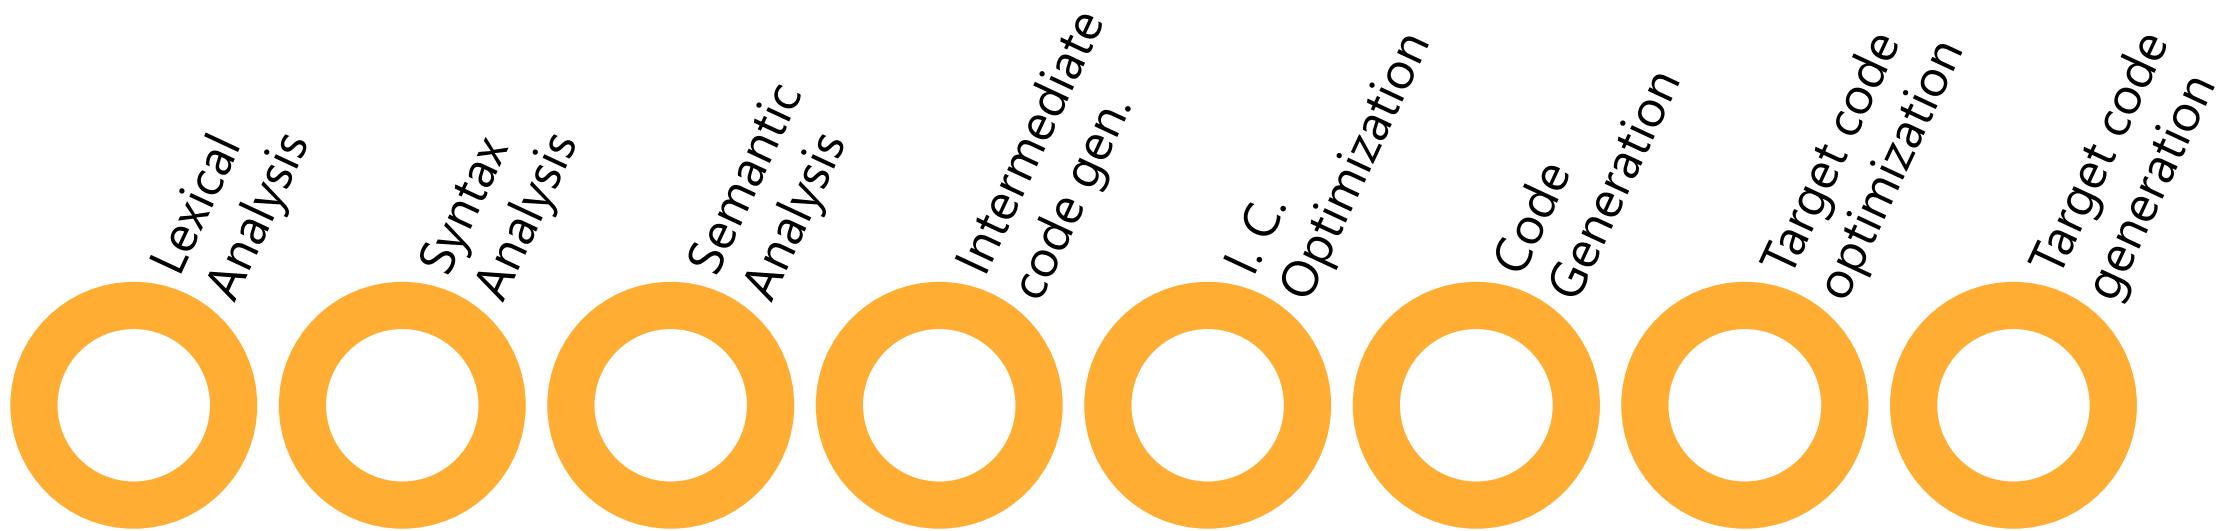
$m \times n$ compilers with $m+n$ components



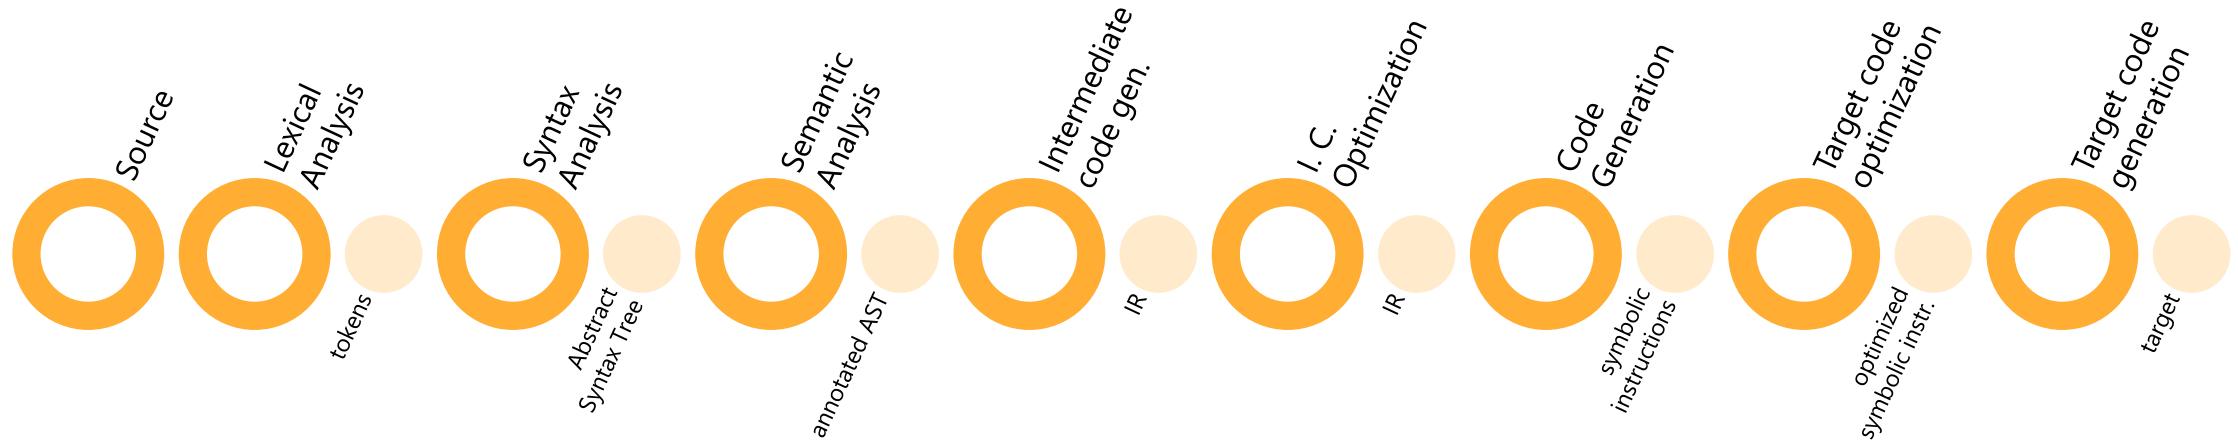
- All language-specific knowledge must be encoded in the front-end
- All target-specific knowledge must be encoded in the back-end



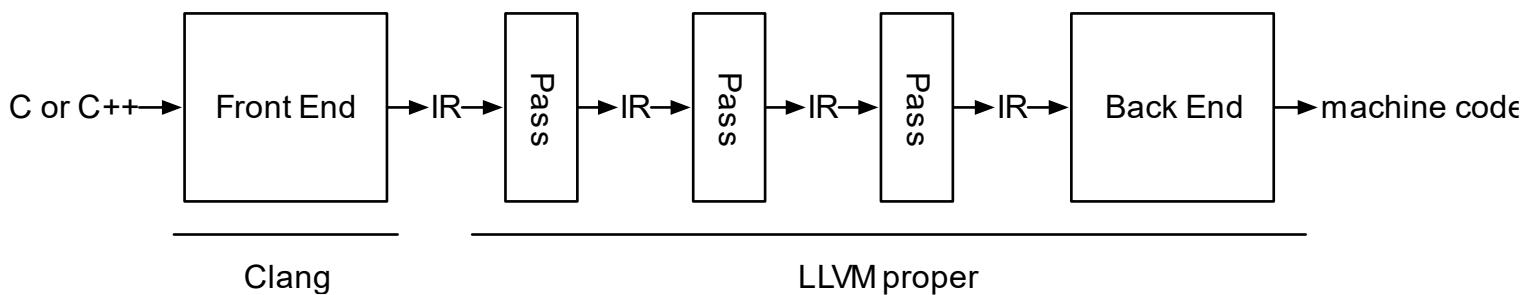
General Compiler Structure



General Compiler Structure



- Example: the **LLVM** compiler



Lexical Analysis

- Reads characters in the source program and groups them into words (basic unit of syntax)
- Produces words and recognizes what sort they are
- The output is called token and is a pair of the form <type, lexeme> or <token_class, attribute>
 - E.g.: **a=b+c** becomes <id, **a**> <=, > <id, **b**> <+, > <id, **c**>
- Needs to record each id attribute: keep a symbol table
 - Lexical analysis eliminates white space, etc
- Speed is important - use a specialized tool
 - e.g., Flex: a tool for generating scanners: programs which recognize lexical patterns in text



Syntax Analysis (Parsing)

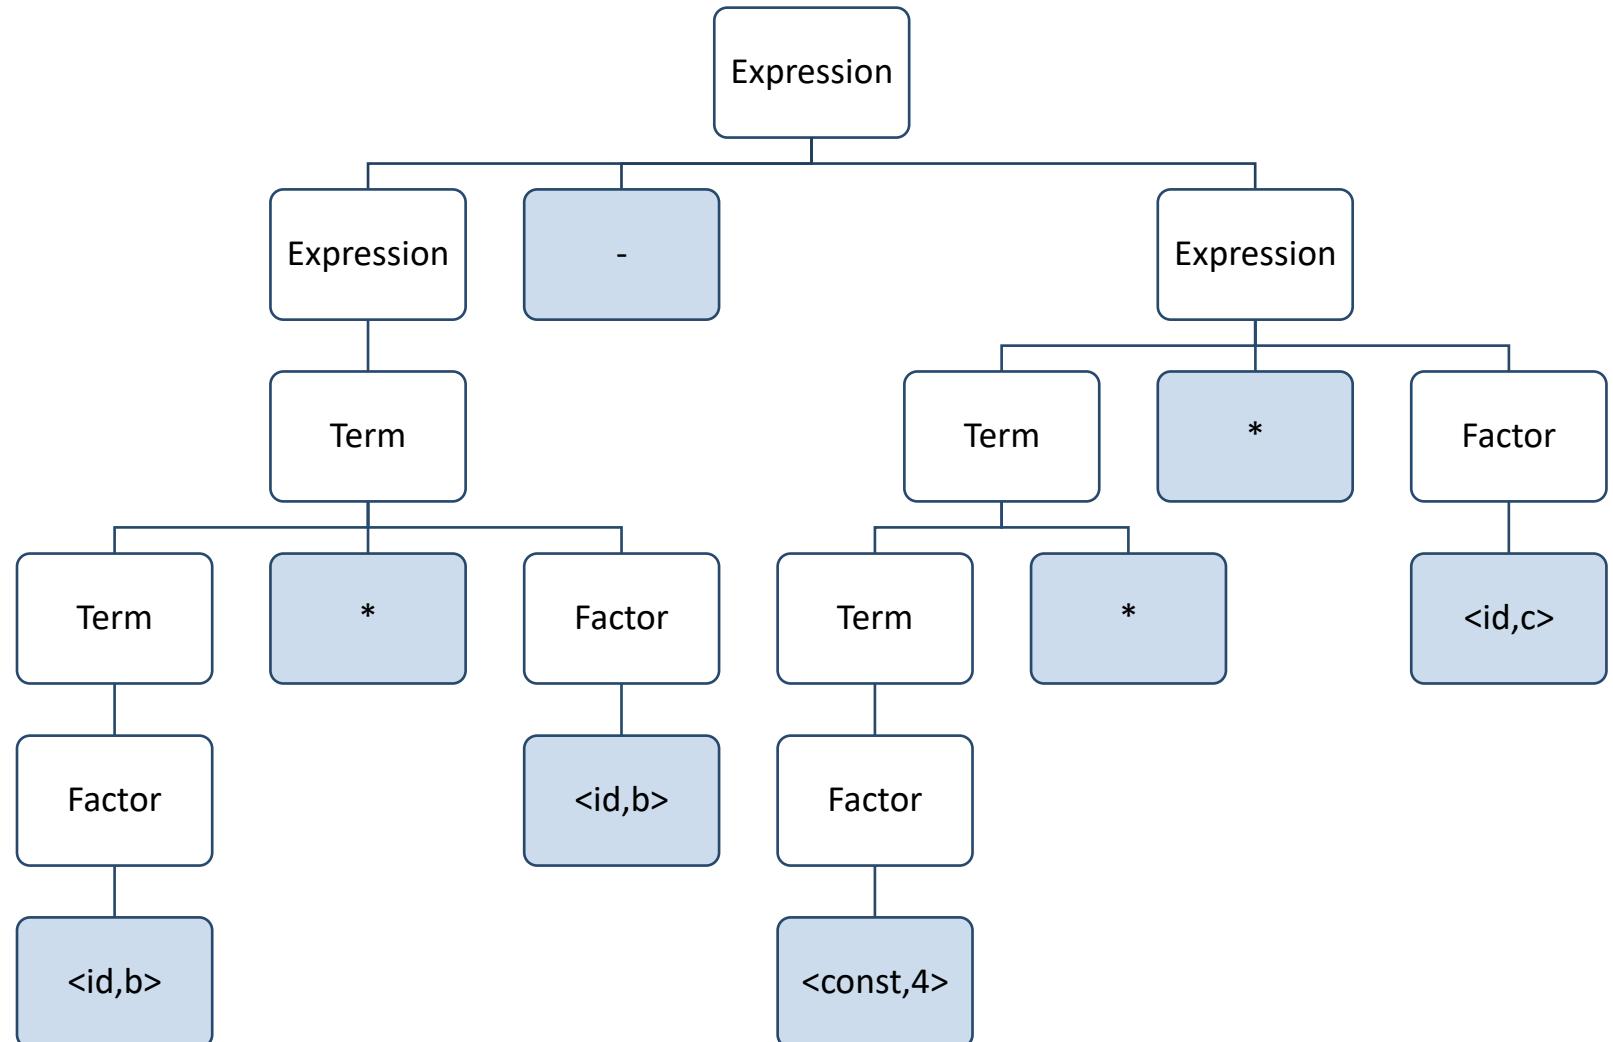
- Imposes a hierarchical structure on the token stream
- This hierarchical structure is usually expressed by recursive rules
- Context-free grammars formalise these recursive rules and guide syntax analysis
- Example
 - A grammar defining simple algebraic expressions

```
expression -> expression '+' term | expression '-' term | term
term -> term '**' factor | term '/' factor | factor
factor -> identifier | constant | '(' expression ')''
```



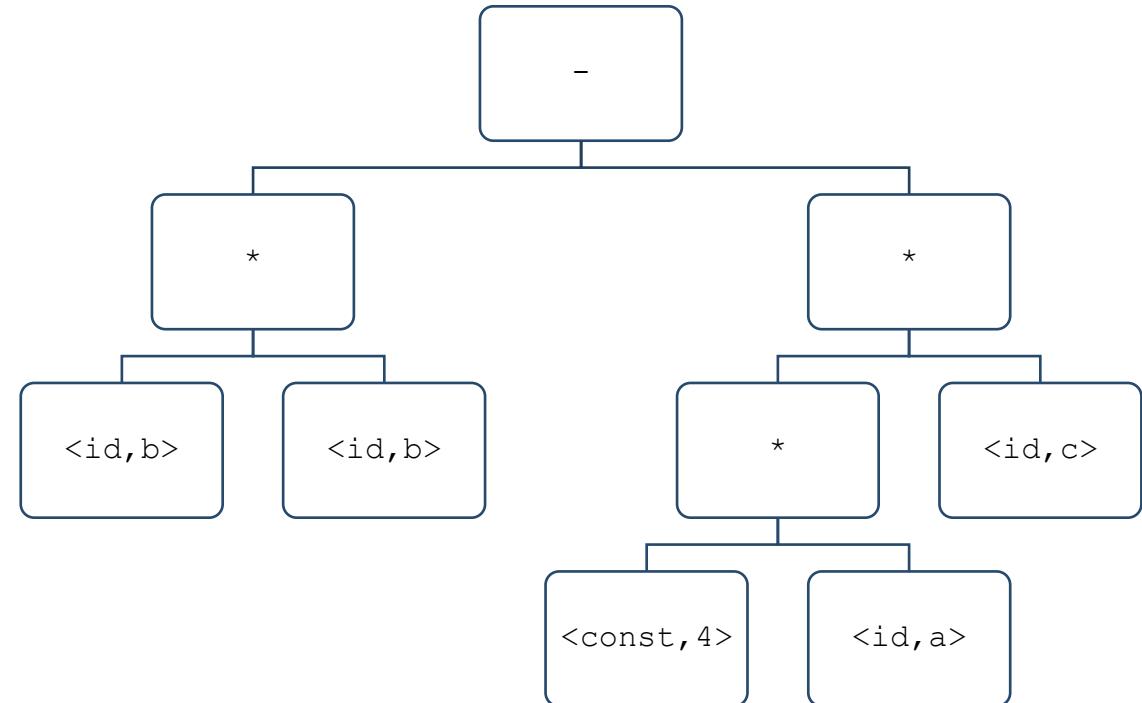
Parsing

- Parse tree for
 $b^*b - 4^*a^*c$



Abstract Syntax Tree (AST)

- Example: an AST for $b^*b - 4^*a^*c$
- An **Abstract Syntax Tree (AST)** is a more useful data structure for internal representation
 - a compressed version of the parse tree
 - summary of grammatical structure without details about its derivation
 - ASTs are a form of IR
- Used as fundamental representation in source-to-source compiler



AST Example: Clang AST

- Example code:

```
int f(int x) {
    int result = (x / 42);
    return result;
}
```

- Compile with:

```
clang -Xclang -ast-dump -fsyntax-only test.cc
```

- Output:

```
TranslationUnitDecl 0x5aea0d0 <<invalid sloc>>
...
`-FunctionDecl 0x5aeab50 <test.cc:1:1, line:4:1> f 'int (int)'
|-ParmVarDecl 0x5aeaa90 <line:1:7, col:11> x 'int'
`-CompoundStmt 0x5aead88 <col:14, line:4:1>
| -DeclStmt 0x5aead10 <line:2:3, col:24>
| `-VarDecl 0x5aeac10 <col:3, col:23> result 'int'
| `-ParenExpr 0x5aeacf0 <col:16, col:23> 'int'
|   `-BinaryOperator 0x5aeacc8 <col:17, col:21> 'int' '/'
|     |-ImplicitCastExpr 0x5aeacb0 <col:17> 'int' <LValueToRValue>
|       `-DeclRefExpr 0x5aeac68 <col:17> 'int' lvalue ParmVar 0x5aeaa90 'x' 'int'
|         `-IntegerLiteral 0x5aeac90 <col:21> 'int' 42
`-ReturnStmt 0x5aead68 <line:3:3, col:10>
  `-ImplicitCastExpr 0x5aead50 <col:10> 'int' <LValueToRValue>
    `DeclRefExpr 0x5aead28 <col:10> 'int' lvalue Var 0x5aeac10 'result' 'int'
```



AST Example: dHPF High Performance Fortran

Fortran code with parallel annotations

```
PROGRAM MAIN
    REAL A(100), X
!HPF$ PROCESSORS P(4)
!HPF$ DISTRIBUTE A(BLOCK) ONTO P
    FORALL (i=1:100) A(i) = X+1
                    CALL FOO(A)
    END

    SUBROUTINE FOO(X)
    REAL X(100)
!HPF$ INHERIT X
    IF (X(1).EQ.0) THEN
        X = 1
    ELSE
        X = X + 1
    END IF
    RETURN
END
```

```
(1 [GLOBAL]
  ((2 [PROG_HEDR]
    (3 [VAR_DECL]
      4 [PROCESSORS_STMT]
      5 [DISTRIBUTE_DECL]
      (6 [FORALL_STMT]
        (7 [ASSIGN_STAT]
          8 [CONTROL_END]
        )
        NULL
      )
      9 [PROC_STAT]
      10 [CONTROL_END]
      ) NULL
    )
    (11 [PROC_HEDR]
      (12 [VAR_DECL]
        13 [INHERIT_DECL]
        (14 [LOGIF_NODE]
          (15 [ASSIGN_NODE]
            16 [CONTROL_END]
          )
          (17 [ASSIGN_NODE]
            18 [CONTROL_END]
          )
        )
        19 [RETURN_STAT]
        20 [CONTROL_END]
      ) NULL
    )
  )
)
```

! 2 components in true-branch
! 6 components in true-branch

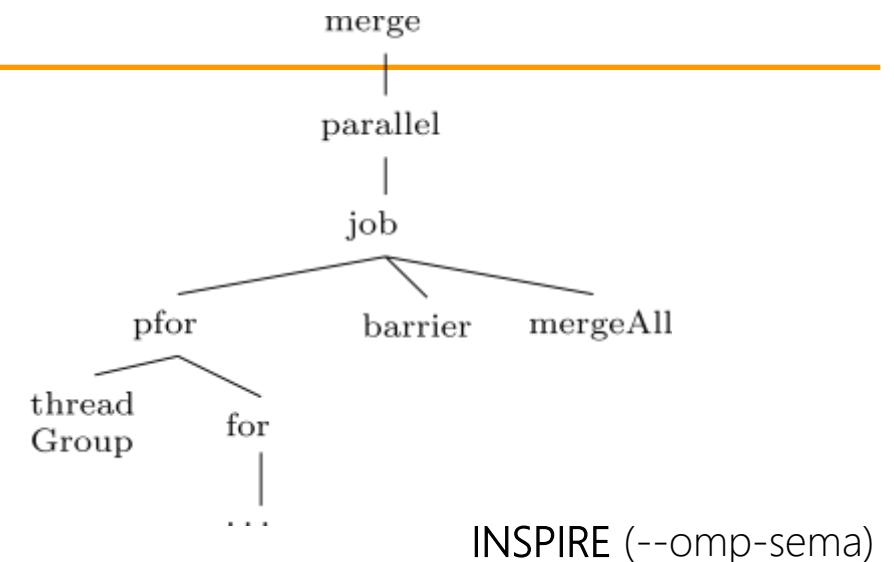
! both branches are non empty



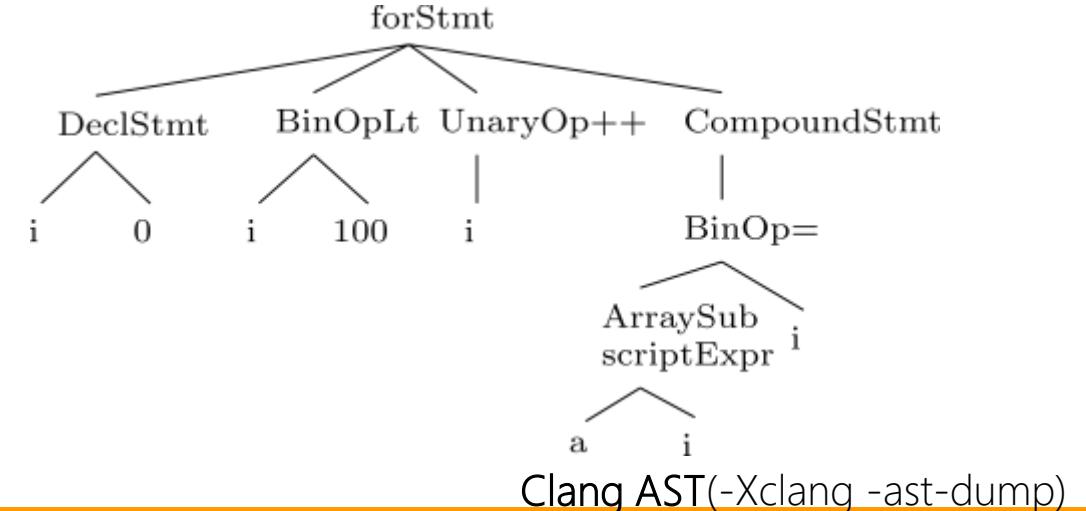
AST Example: the Insieme Parallel Compiler

- Source-to-source compiler, parallel by design
- INSPIRE (INSieme Parallel Intermediate Representation)
 - high-level IR
 - all transformations happens at AST level
 - output is C/C++ code enabling parallelization with pthreads, MPI, OpenCL
 - parallelism is expressed in the IR
- Example

```
int a [N];
#pragma omp parallel
{
    #pragma omp for
    for(int i=0; i<N; i++) {
        a[i]=i;
    }
}
```



INSPIRE (--omp-sema)



Clang AST(-Xclang -ast-dump)



Semantic Analysis (Context Handling)

- Collects context (semantic) information, checks for semantic errors, and annotates nodes of the tree with the results
- Examples
 - type checking: report error if an operator is applied to an incompatible operand
 - check flow-of-controls
 - uniqueness or name-related checks



Intermediate Code Generation

- Translate language-specific constructs in the AST into more general constructs
- Example of a form of IR (3-address code)

```
tmp1 = 4
tmp2 = tmp1*a
tmp3 = tmp2*c
tmp4 = b*b
tmp5 = tmp4-tmp3
```

- Typically organized in a control flow graph



Control Flow Graph

- **Control-flow graph** (CFG): a graph-based representation of all paths that might be traversed through a program during its execution
 - each node is an instruction or a **basic block**
 - edges represent jump in the control flow
- **Basic block**
 - straight-line code sequence with no branches in except to the entry and no branches out except at the exit
 - only one entry block, only one exit block



Code Optimization

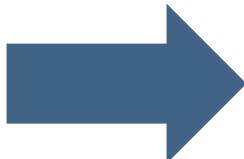
- The goal is to improve the intermediate code and, thus, the effectiveness of code generation and the performance of the target code
- Optimizations can range from trivial (e.g., constant folding) to highly sophisticated (e.g., in-lining)
 - Example: replace the first two statements in the example of the previous slide with:
`tmp2=4*a`
- Modern compilers perform such a range of optimizations, that one could argue for:



Optimizations

- Example: Dead Code Elimination (**DCE**)

```
int global;
void f ()
{
    int i;
    i = 1;
    global = 1;
    global = 2;
    return;
    global = 3;
}
```



```
int global;
void f ()
{
    global = 2;
    return;
}
```



Code Generation

- Map the IR onto a linear list of target machine instructions in a symbolic form
- Three major steps
 1. **instruction selection**: a pattern matching problem
 2. **register allocation**: each value should be in a register when it is used (but there is only a limited number): NP-Complete problem
 3. **instruction scheduling**: take advantage of multiple functional units: NP-Complete problem
- Target, machine-specific properties may be used to optimize the code
- Finally, machine code and associated information required by the Operating System are generated



Other Representations used in Compilers

- Call graph
 - shows dependences between procedures
 - useful for inter-procedural analysis and optimization
 - E.g., inlining using Strongly Connected Component (Tarjan's algorithm)
- Data dependence graph
 - encodes the flow of data
 - node: program statement
 - edge: connects two nodes if one uses the result of the other
 - useful in examining the legality of program transformations

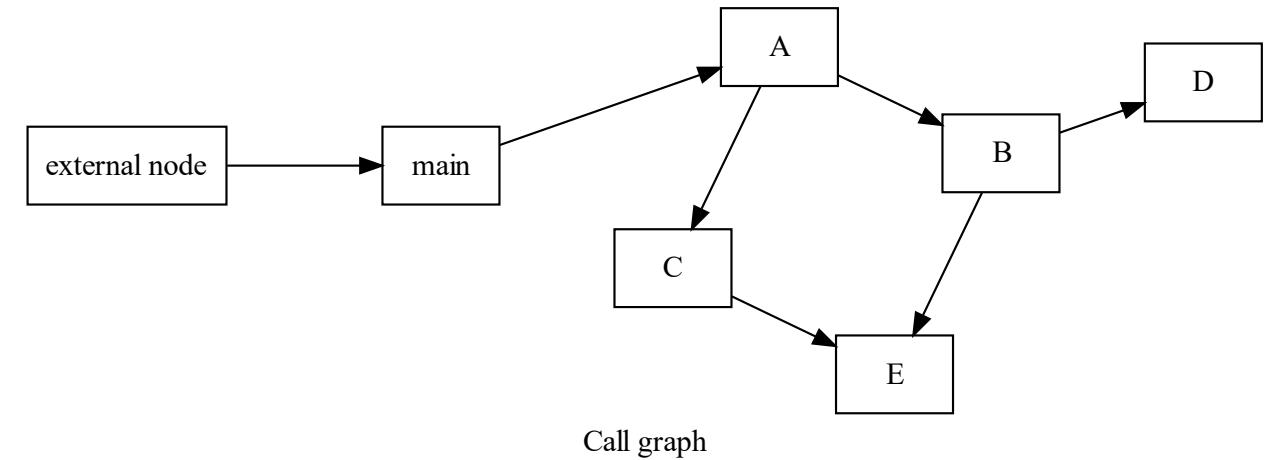


Example: Call Graph in Clang/LLVM

- Example code:

```
static void E() { }
static void D() { }
static void C() { E(); }
static void B() { D(); E(); }
static void A() { B(); C(); }
int main() {
    A();
}
```

- Output (circo svg engine):



- Compile: **clang -S -emit-llvm main.c -o - | opt -analyze -dot-callgraph**
- Visualize: **dot -Tpng -ocallgraph.png callgraph.dot**



Getting Started with LLVM

- You can generate LLVM IR (also called bitcode) with the following command

```
clang example.c -o example.bc -c -emit-llvm
```

- you can quickly run it by using the LLVM interpreter

```
lli example.bc
```

- The bitcode can be browsed by using the LLVM disassembler

```
llvm-dis < example.bc
```

- To produce assembly from the bytecode with the LLVM backend compiler

```
llc example.bc -o example.s
```



LLVM Basic Block Example

■ Code example

```
void test(bool b) {  
    int a;  
    if (b) {  
        a = 1;  
    }  
    printf("%d\n", a);  
}
```

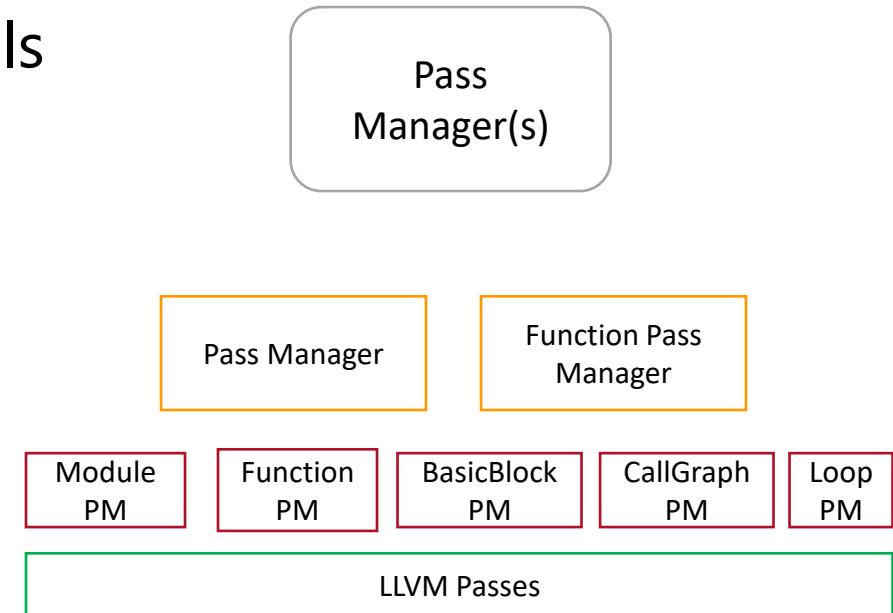
■ LLVM bitcode

```
define void @test(i1 zeroext) #0 !dbg !6 {  
    %2 = alloca i8, align 1  
    %3 = alloca i32, align 4  
    %4 = zext i1 %0 to i8  
    store i8 %4, i8* %2, align 1  
    call void @llvm.dbg.declare(metadata i8* %2, metadata !10, metadata !11), !dbg !12  
    call void @llvm.dbg.declare(metadata i32* %3, metadata !13, metadata !11), !dbg !15  
    %5 = load i8, i8* %2, align 1, !dbg !16  
    %6 = trunc i8 %5 to i1, !dbg !16  
    br i1 %6, label %7, label %8, !dbg !18  
  
; <label>:7:                                     ; preds = %1  
    store i32 1, i32* %3, align 4, !dbg !19  
    br label %8, !dbg !21  
  
; <label>:8:                                     ; preds = %7, %1  
    %9 = load i32, i32* %3, align 4, !dbg !22  
    %10 = call i32 (i8*, ...) @printf(...), !dbg !23  
    ret void, !dbg !24  
}
```



The LLVM Compilation Infrastructure

- We can build compiler passes at different levels
 - Basic block
 - Function
 - Loop
 - Region
 - Call graph



LLVM BasicBlock Pass

- An example BB pass that prints all instructions

```
#include <llvm/IR/Type.h>
#include <llvm/Pass.h>
#include <llvm/IR/BasicBlock.h>
#include <llvm/Support/raw_ostream.h>
using namespace llvm;
namespace {
class MyBlockPass : public BasicBlockPass {
public:
    static char ID;
    MyBlockPass() : BasicBlockPass(ID) {}

    virtual bool runOnBasicBlock(BasicBlock &BB) {
        errs().write_escaped(BB.getName()) << '\n';
        // iterating instructions in the current BasickBlock
        for(Instruction &i : BB) {
            errs() << " - " << i.getOpcodeName() << " ";
            Type *type = i.getType();
            type->print(errs());
            errs() << '\n';
        }
        errs() << '\n';
        return false;
    }
};

} // namespace

char MyBlockPass ::ID = 0;
static RegisterPass<MyBlockPass> X("block-pass", "Block Pass");
```



LLVM Function Pass

```
#include <llvm/Pass.h>
#include <llvm/IR/Function.h>
#include <llvm/Support/raw_ostream.h>
using namespace llvm;
namespace {
class MyFunctionPass : public FunctionPass {
public:
    static char ID;
    MyFunctionPass () : FunctionPass(ID) {}
    virtual bool runOnFunction(Function &F) {
        errs().write_escaped(F.getName()) << '\n';
        // iterate arguments
        errs() << " - args:" << '\n';
        for(Argument &a : F.getArgumentList()) {
            errs() << "    - ";
            errs().write_escaped(a.getName()) << '\n';
        }
        // iterate BB in a function
        errs() << " - blocks:" << '\n';
        for(BasicBlock &b : F.getBasicBlockList()) {
            errs() << "      * ";
            errs().write_escaped(b.getName()) << '\n';
        }
        errs() << '\n';
        return false;
    }
};
} // namespace
char MyFunctionPass::ID = 0;
static RegisterPass<MyFunctionPass> X("function-pass", "FunctionPas
```



LLVM Loop Pass

- Warning: there is not explicit loop in a CGF
- Only works after Loop Analysis pass
 - Identifies natural loop

```
#include <llvm/Analysis/LoopPass.h>
#include <llvm/Support/raw_ostream.h>
using namespace llvm;

namespace {

class MyLoopPass : public LoopPass {
public:
    static char ID;
    MyLoopPass() : LoopPass(ID) {}

    bool runOnLoop(Loop *L, LPPassManager &LPPM) {
        L->print(errs());
        for(BasicBlock *b : L->getBlocks()) {
            errs() << " - ";
            errs().write_escaped(b->getName()) << '\n';
        }
        errs() << '\n';
        return false;
    }
};

// namespace

char MyLoopPass::ID = 0;
static RegisterPass<MyLoopPass> X("loop-pass", "Loop Pass");
```



What Compilation Passes are Executed with Clang/LLVM?

- Basic passes (no opt)

```
clang -mllvm -opt-bisect-limit=-1 test.cc
```

- O3 passes

```
clang -O3 -mllvm -opt-bisect-limit=-1 test.cc
```



Lab

- Experiment with LLVM
 - Generate AST and LLVM bitcode code for C/C++ code
 - Compare LLVM bitcode with different compilation flags



Instruction Scheduling, Compiler-based Optimizations

High Performance Computing, Summer 2021



Biagio Cosenza
Department of Computer Science
University of Salerno
bcosenza@unisa.it

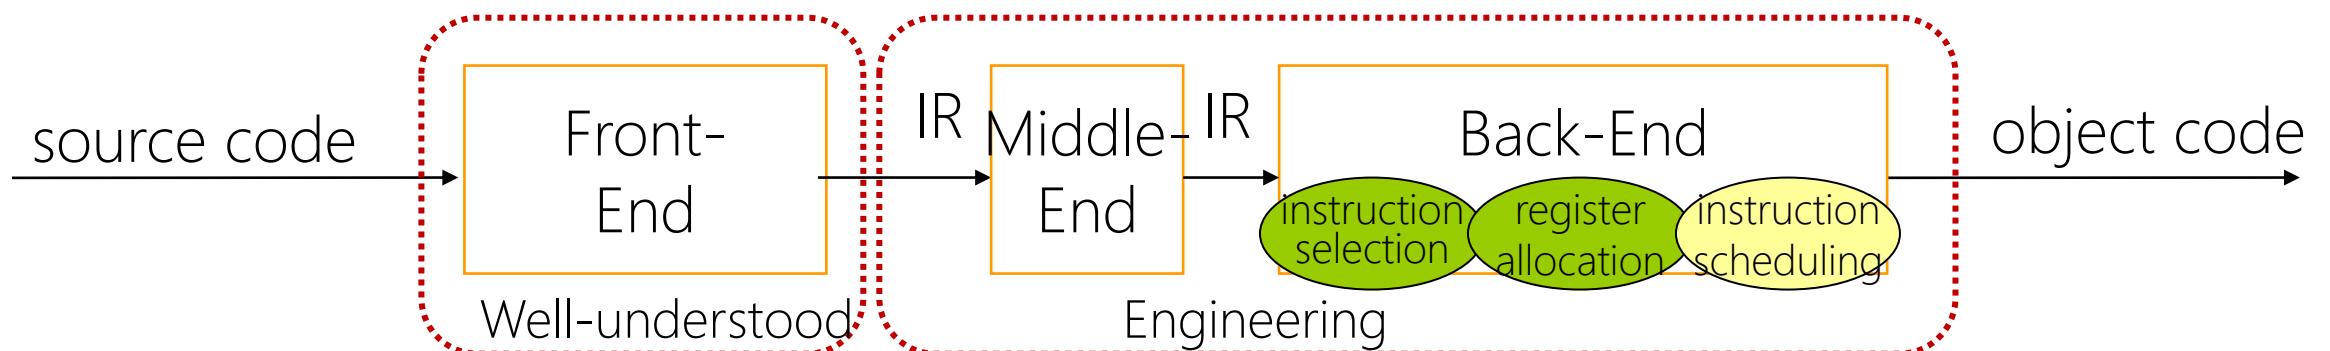
Outline

- Instruction scheduling
- Optimization overview
 - Loop optimizations



Instruction Scheduling

- The problem
 - given a code fragment for some target machine and the latencies for each individual operation, reorder operations to minimize execution time
 - recall: modern processors may have multiple functional units
- The task
 - produce correct code; minimize wasted cycles; avoid spilling registers; operate efficiently



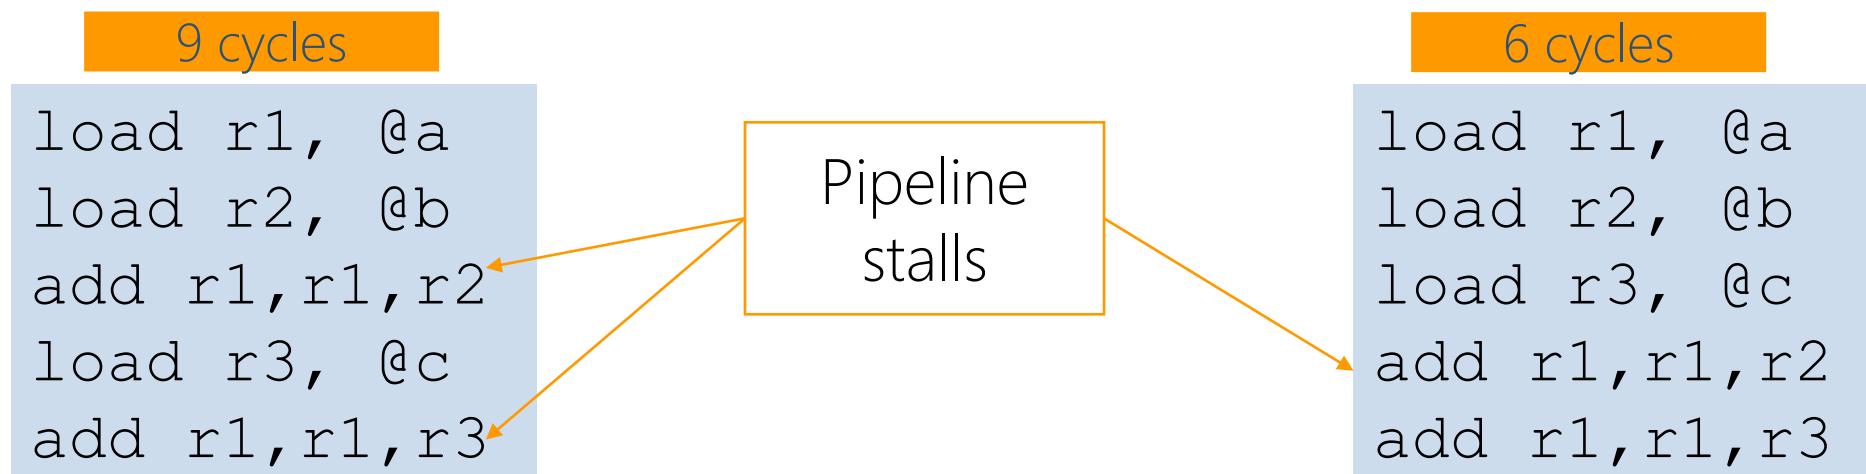
Instruction Scheduling: Background

- Many operations have delay latencies for execution
 - E.g., load, store: <delay> CPU cycles (depends on the processor)
 - issue load, result appears <delay> cycles later
 - execution continues unless result is referenced
 - premature reference causes hardware to stall
- Modern machines can issue several operations per cycle (**pipelining**)
- Execution time is order-dependent (has been since the 60s)
- Overview of a solution
 - move loads back at least <delay> slots from where they are needed, but this increases register pressure (i.e., more registers may be needed)
 - ideally, we want to minimize both hardware stalls and added register pressure



Motivating Example

- Two variants to compute $(a+b)+c$
 - assume that the latency of a load is 3 cycles; all other instructions have a latency of 1 cycle
 - just an example with simplified hardware model
 - note: costs due to the memory hierarchy are not taken in account



Instruction Scheduling Algorithm for a Basic Block

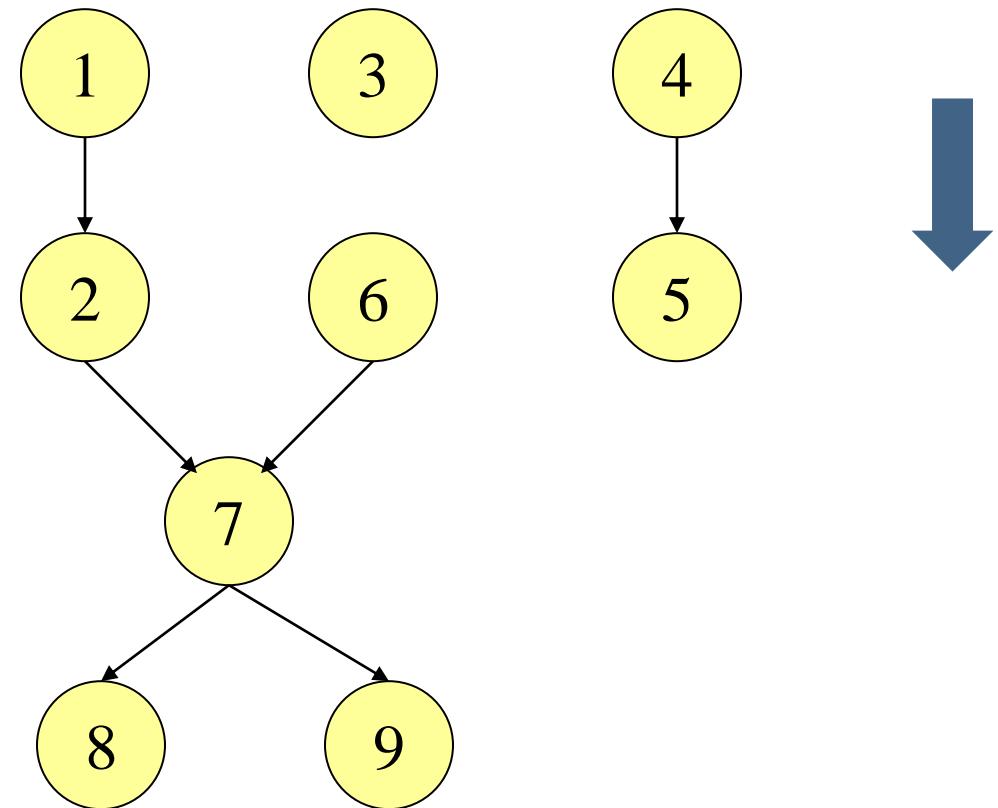
1. Build a **precedence** (data dependence) graph
2. Compute a **priority function** for the nodes of the graph
3. Use **list scheduling** to construct a schedule, one cycle at a time
 1. use a queue of operations that are ready
 2. at each cycle
 - choose a ready operation and schedule it
 - update the ready queue
 - A greedy heuristic; open to variations
 - recall: greedy heuristic: an algorithmic technique in which an optimization problem is solved by finding locally optimal solutions



1. Build a Precedence Graph

- Dependence graph for the following code

```
1. load r1, @x
2. load r2, [r1+4]
3. and r3, r3, 0x00FF
4. mult r6, r6, 100
5. store r6
6. div r5, r5, 100
7. add r4, r2, r5
8. mult r5, r2, r4
9. store r4
```



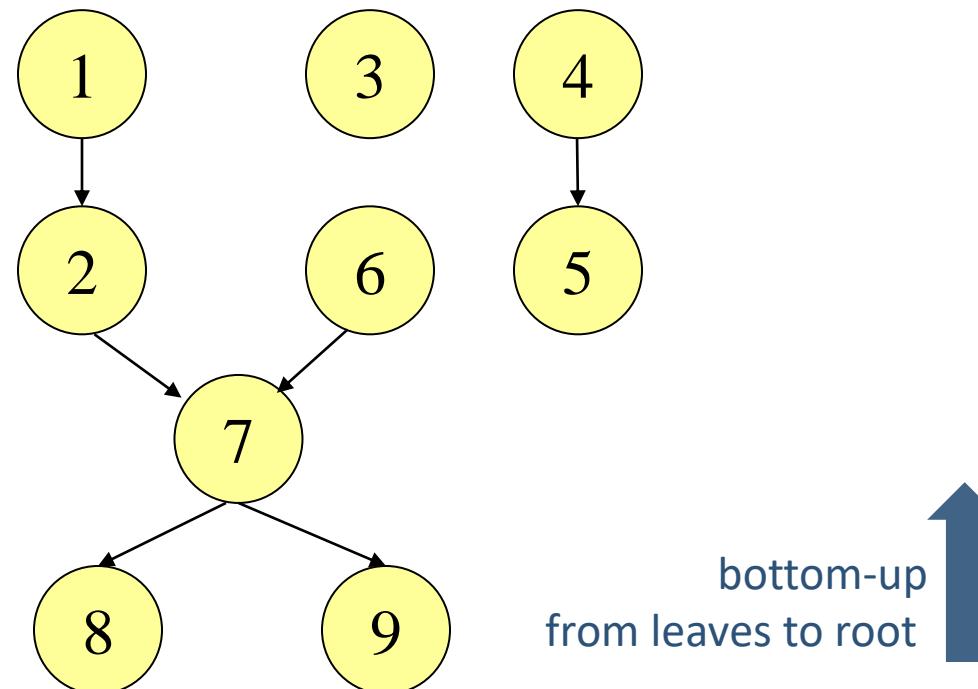
2. Compute a Priority Function

- Assign to each node a weight equal to the longest delay latency (total) to reach a leaf in the graph from this node (include latency of current node)

$$\text{weight}_i = \text{latency}_i + \max(\text{weight}_{\text{all successor nodes}})$$

- Assume
 - div 4 cycles
 - mult 3 cycles
 - 1 cycle for the rest

```
1. load r1, @x
2. load r2, [r1+4]
3. and r3, r3, 0x00FF
4. mult r6, r6, 100
5. store r6
6. div r5, r5, 100
7. add r4, r2, r5
8. mult r5, r2, r4
9. store r4
```



node	weight
1	6
2	5
3	1
4	4
5	1
6	8
7	4
8	3
9	1



3. Local List Scheduling Algorithm

```
Cycle=1; Ready = set of available operations; Active = {}
```

```
while (Ready  $\cup$  Active  $\neq \{ \}$ )
```

```
    if (Ready  $\neq \{ \}$ ) then
```

```
        remove an op from Ready (based on the weight)
```

```
        schedule(op) = cycle
```

```
        Active = Active  $\cup$  op
```

```
    cycle = cycle+1
```

```
    for each op in Active
```

```
        if (schedule(op) + delay(op)  $\leq$  cycle) then
```

```
            remove op from Active
```

```
            for each immediate successor s of op
```

```
                if (all operand of s are available) then
```

```
                    Ready = Ready  $\cup$  s
```

Removal in priority order

if op has completed execution

if all successor's op are ready, then put it on ready

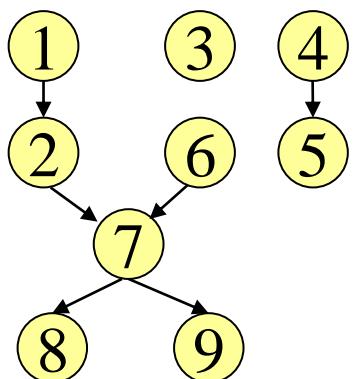


Finding the Schedule

- Final schedule

Cycle	Instructions ready	Schedule	Instructions active
1	6, 1, 4, 3	div r5, r5, 100	6
2	1, 4, 3	load r1, @x	6, 1
3	2, 4, 3	load r2, [r1+4]	6, 2
4	4, 3	mult r6, r6, 100	6, 4
5	7, 3	add r4, r2, r5	4, 7
6	8, 3, 9	mult r5, r2, r4	4, 8
7	3, 9, 5	and r3, r3, 0x00ff	8, 3
8	9, 5	store r4	8, 9
9	5	store r6	5

node	weight
1	6
2	5
3	1
4	4
5	1
6	8
7	4
8	3
9	1



More on Scheduling

- Two distinct classes of list scheduling
 - **forward list scheduling**: start with all available operations; work forward in time (Ready: all operands available)
 - **backward list scheduling**: start with leaves; work backward in time (Ready: latency covers uses)
 - folk wisdom is to try both and keep the best result
- Variations on computing priority function
 - maximum path length containing node (decreases register usage)
 - prioritize critical path
 - number of immediate successors or total number of descendants
 - increment weight if node contains a last use (shortens live ranges)
 - do not add latency to node's weight
 - maximum delay latency from first available node



Example: Extend List Scheduling for Multiple Functional Units

- Modern architectures can run operations in parallel
- List scheduling needs to be modified so that it can schedule as many operations per cycle as functional units (assuming that there are instructions available)

- 3rd line in the list scheduling algorithm changes to:

```
while (Ready != {} &&
      there_are_free_functional_units)
```

- Back to the previous example

- assume two functional units that can issue any instruction

F. U. 1	F. U. 2	Ready Set
6. div r5,r5,100	1. load r1,@x	{6,1,4,3}
2. load r2,[r1+4]	4. mult r6,r6,100	{2,4,3}
3. and r3,r3,0x00FF	nop	{3}
nop	nop	{ }
7. add r4,r2,r5	5. store r6	{7,5}
8. mult r5,r2,r4	9. store r4	{8,9}



Exercise

- Assume two functional units: one for ALU operations only and another for memory operations only. A load and a mult have a latency of 2 cycles, all other instructions have a latency of 1 cycle

```
1. load r1, @x
2. load r2, @y
3. add r2, r2, 42
4. load r3, @z
5. shl r4, r1, 4
6. store r4
7. mult r5, r2, r3
8. add r6, r5, r4
9. store r5
```

ALU	MEM
nop	2
nop	4
3	1
7	nop
5	nop
8	9
nop	6



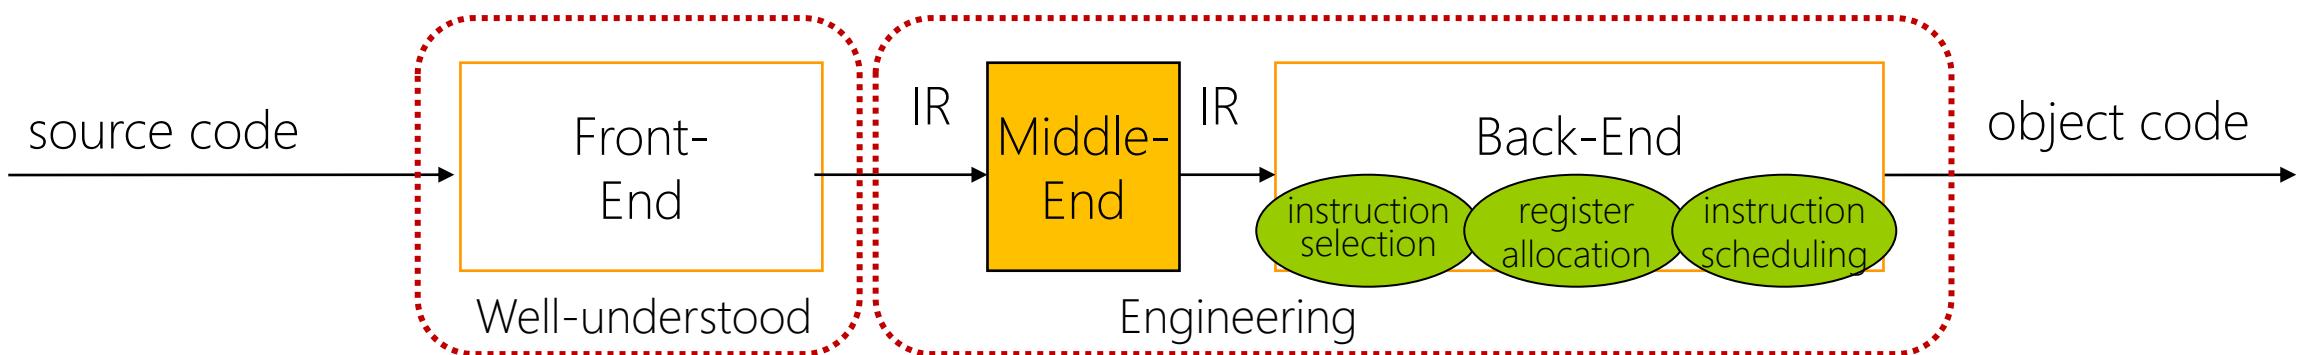
Further Issues

- Going beyond basic blocks
 - identify high-frequency path and schedule as if a single block
- Modulo scheduling
 - schedule multiple iterations together
 - run several iterations concurrently - i.e., overlap successive iterations
- Register allocation with instruction scheduling
 - the former before the latter restricts the choices for scheduling
 - the latter before the former: if register allocation has to spill registers, the whole (carefully done) schedule changes!
- Besides performance, we may want to minimize power consumption, size of the code, ...



Code Optimization

- Goal
 - improve program performance within some constraints (may also reduce size of the code, power consumption, etc...)
- Issues
 - **legality**: must preserve the meaning of the program
 - externally observable meaning may be sufficient/may need flexibility
 - **benefit**: must improve performance on average or common cases
 - predicting program performance is often non-trivial
 - **compile-time cost justified**: list of possible optimizations is huge
 - inter-procedural optimizations



“Optimizing” Transformations

- Finding an appropriate sequence of transformations is a major challenge: modern optimizers are structured as a series of passes
 - optimization 1 is followed by optimization 2; optimization 2 is followed by optimization 3, and so on...
- Transformations may improve program at
 - source level (algorithm specifics)
 - IR (machine-independent transformations)
 - target code (machine-dependent transformations)
- Some typical transformations
 - discover and propagate some constant value
 - remove unreachable/redundant computations
 - encode a computation in some particularly efficient form



A Classification of Compiler Optimizations

- By Scope
 - **local**: within a single basic block
 - **peephole**: on a window of instructions (usually local)
 - **loop-level**: on one or more loops or loop nests
 - **global**: for an entire procedure
 - **inter-procedural**: across multiple procedures or whole program (also called IPO)
- Example: LLVM Passes
 - Module, CallGraph, Function, Loop, Region, BasicBlock
- By machine information used
 - machine-independent versus machine-dependent
- By effect on program structure
 - algebraic transformations
 - $x+0, x*1, 3*z^4, \dots$
 - affine transformations on polyhedral loop transf.
 - reordering transformations
 - change the order of two computations
 - loop-level reordering transformations



Transformations (1)

- Common Subexpression Elimination

```
A[i][i*2+10] = B[i][i*2+10]+5;
```



```
tmp = i*2+10;  
A[i][tmp] = B[i][tmp]+5;
```

- Copy Propagation

```
t = i*4;  
s = t;  
a[s] = a[s]+4;
```



```
t = i*4;  
a[t] = a[t]+4;
```

- Constant Propagation

```
N=64  
c=2  
for (i=0;i<N;i++)  
    a[i] = a[i]+c;
```



```
N=64  
c=2  
for (i=0;i<64;i++)  
    a[i] = a[i]+2;
```



Transformations (2)

- Constant folding

```
tmp = 5*3 + 8 - 12/2;
```



```
tmp = 17;
```

- Dead Code Elimination

```
if(3 > 7) {  
    x = a + b;  
}
```



```
// removed
```

- Reduction in strength

```
x*2 + x*1024;
```



```
x + x + (x<<10);
```



Loop Transformations

- Very important for performance
 - change the order in which the iteration space is traversed
 - can expose parallelism; increase available ILP; improve memory behavior
- Dependence testing is required to check validity of transformation
- Automatic parallelization approaches are mainly based on loop-parallelism



Loop Merge

```
for (ia=exp1; ia<exp2; ia++)  
    A (ia) ;  
for (ib=exp1; ib<exp2; ib++)  
    B (ib) ;
```



```
for (i=exp1; i<exp2; i++) {  
    A (i) ;  
    B (i) ;  
}
```

- Improve locality
- Reduce loop overhead



Loop Merge Example

```
for (i=0; i<N; i++)  
    B[i] = f(A[i]);  
for (j=0; j<N; j++)  
    C[j] = f(B[j],A[j]);
```



```
for (i=0; i<N; i++) {  
    B[i] = f(A[i]);  
    C[i] = f(B[i],A[i]);  
}
```

- Dependency check
- Different memory usage, may improve locality
 - Is it always better?



Loop Body Split

```
for (i=exp1; i<exp2; i++) {  
    A(i) ;  
    B(i) ;  
}
```



```
for (ia=exp1; ia<exp2; ia++)  
    A(ia) ;  
for (ib=exp1; ib<exp2; ib++)  
    B(ib) ;
```



Example: Loop Body Split + Merge

```
for (i=0; i<N; i++) {  
    A[i] = f(A[i-1]);  
    B[i] = g(in[i]);  
}  
for (j=0; j<N; j++)  
    C[i] = h(B[j],A[N]);
```

```
for (i=0; i<N; i++)  
    A[i] = f(A[i-1]);  
    for (j=0; j<N; j++) {  
        B[j] = g(in[j]);  
        C[j] = h(B[j],A[N]);  
    }
```



```
for (i=0; i<N; i++)  
    A[i] = f(A[i-1]);  
    for (k=0; k<N; k++)  
        B[k] = g(in[k]);  
    for (j=0; j<N; j++)  
        C[j] = h(B[j],A[N]);
```



Loop Interchange

```
for (ia=exp1; ia<exp2; ia++) {  
    for (ib=exp3; ib<exp4; ib++) {  
        A(ia, ib);
```



```
for (ib=exp3; ib<exp4; ib++) {  
    for (ia=exp1; ia<exp2; ia++) {  
        A(ia, ib);
```

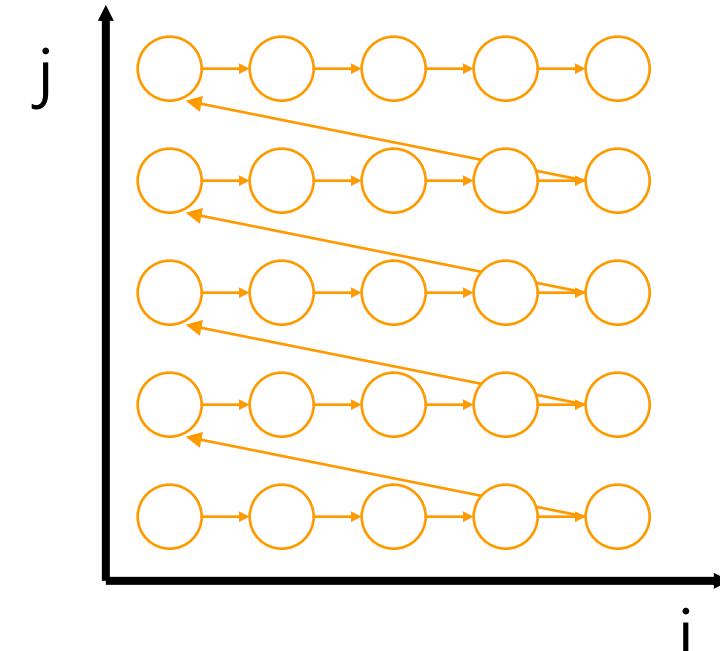
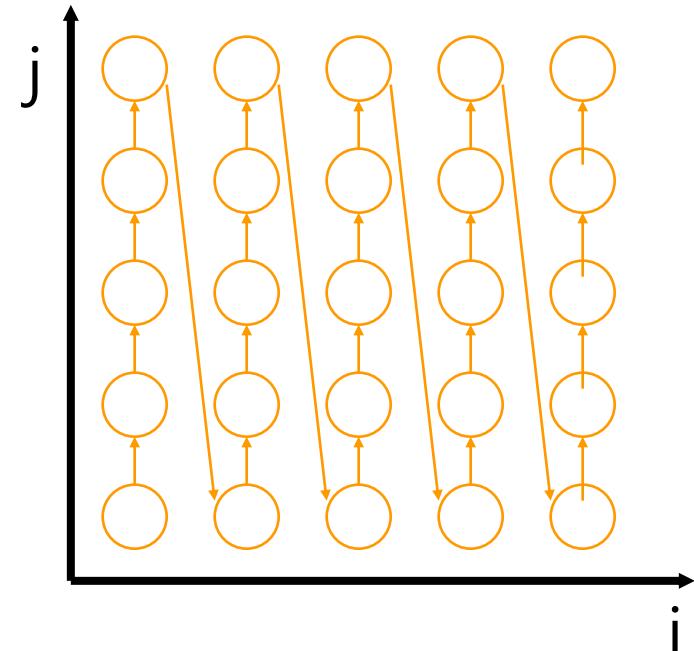


Loop Interchange Example

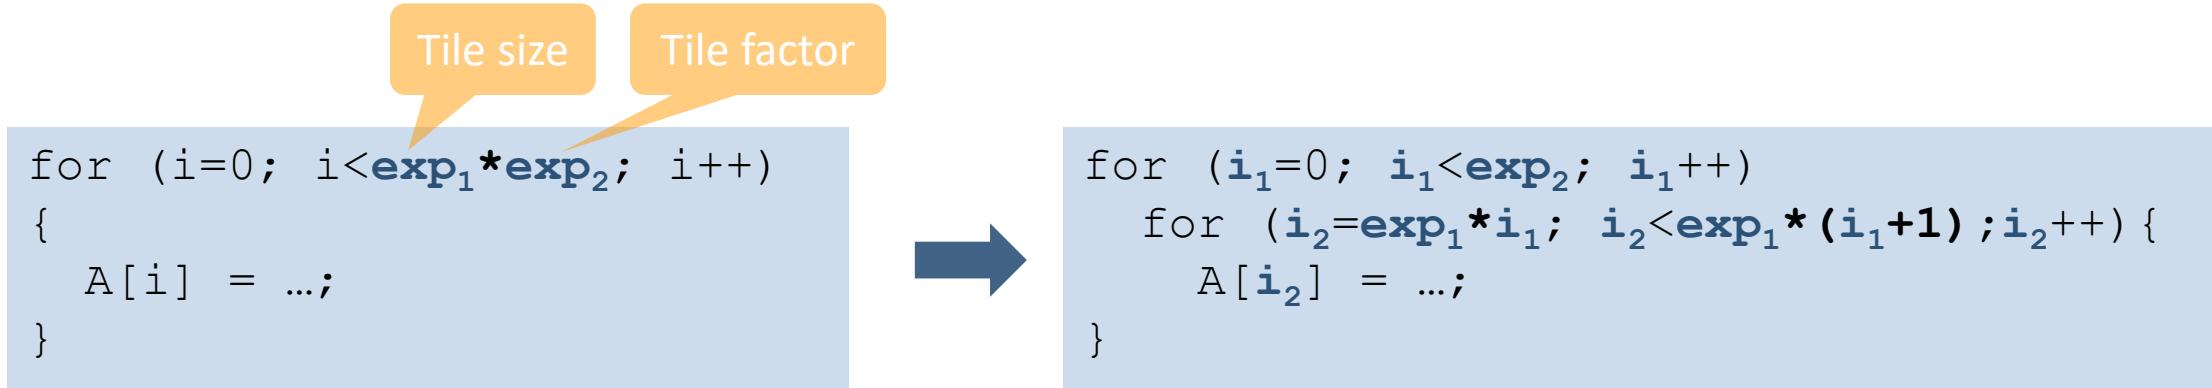
```
for(j=0; j<H; j++)  
    for(i=0; i<W; i++)  
        A[i][j] = ...;
```



```
for(i=0; i<W; i++)  
    for(j=0; j<H; j++)  
        A[i][j] = ...;
```



Loop Tiling



- Improve cache reuse by dividing the iteration space into tiles and iterating over these tiles
- Useful when the working set does not fit into cache or when there exists much interference
 - good with parallel loop (e.g., multiple threads can work on the same data)
- Two adjacent loops can legally be tiled if they can legally be interchanged



2D Tiling Example

```
for(i=0; i<N; i++)  
    for(j=0; j<N; j++)  
        A[i][j] = B[j][i];
```



```
for(TI=0; TI<N; TI+=16)  
    for(TJ=0; TJ<N; TJ+=16)  
        for(i=TI; i<min(TI+16,N); i++)  
            for(j=TJ; j<min(TJ+16,N); j++)  
                A[i][j] = B[j][i];
```

- Two-dimensional case
 - from 2 to 4 loops
 - it can be further extended to N-dimensional cases
- Most tile size selection algorithms use a cache model
 - generate collection of tile sizes
 - estimate resulting cache miss rate



Loop Unrolling

```
for (i=exp1; i<exp2; i++)  
    A[i]=...;
```



```
for (i=exp1/2; i<exp2/2; i++) {  
    A[2*i]=...;  
    A[2*i+1]=...;  
}
```

- Duplicate loop body and adjust loop header
- Increases available ILP, reduces loop overhead, and increases possibilities for common subexpression elimination
- Always valid



Peephole Transformation

- Optimization performed over a very small set of instructions in a segment of generated code
 - the set is called a **peephole** or a **window**
 - recognize sets of instructions that can be replaced by shorter or faster sets of instructions (replacement rules)
- Examples
 - constant folding: evaluate constant sub-expression in advance
 - strength reduction: replace slow operations with faster equivalents
 - null sequences: delete useless operations
 - combine operations: replace several operations with one equivalent
 - algebraic laws: use algebraic laws to simplify or reorder instructions
 - special case instructions: use instructions designed for special operand cases
 - address mode operations: use address modes to simplify code



Binary Translation (or Binary Recompilation)

- **Binary translation** is the emulation of one instruction set by another through translation of binary code
 - sequences of instructions are translated from the source to the target instruction set
 - the target instruction set may be the same as the source instruction set, providing testing and debugging features such as instruction trace, conditional breakpoints and hot spot detection
- may be implemented
 - as **peephole** transformation
 - with **instruction lifting**
 - e.g., from x86 to LLVM IR (bitcode)
- **Static** binary translation
 - difficult: not all code can be discovered by a translator, e.g., indirect branches (value known at runtime)
- **Dynamic** binary translation
 - apply to a short sequence of code (e.g., basic block)
 - when possible, and branch instructions are made to point to already translated and saved code (**memoization**)



Binary Translation: Examples

- Static
 - x86 to ARM translator
 - x86 to x64
- Dynamic
 - QuickTransit (Apple/Transitive): SPARC to x86, x86 to Power architecture, MIPS to Itanium, PPC to x86
 - Intel : from IA-32 to Itanium
- Hardware binary translation
 - x86 Intel CPUs since the Pentium Pro translate complex CISC x86 instructions to more RISC-like internal micro-operations

